

Cosma

Omnilingual Semantic Search Platform

Technical Whitepaper - Investor Edition

Azetta.ai

Version 1.0 - December 27, 2025

Confidential & Proprietary

Do not distribute without permission.

Executive Summary

Azetta.ai delivers **Cosma**, an enterprise-grade omnilingual semantic search platform built on proprietary vector database technology and unified multilingual embeddings. Our platform enables organizations to search across 100+ languages with a single query, deployed as a fully managed SaaS solution.

Contents

1	The Team	5
1.1	Founding Team	5
1.1.1	Taha Bouhsine - Co-Founder & CEO	5
1.1.2	Poppy Yuan - Co-Founder & COO	6
1.1.3	Douglas Seo - Co-Founder & CTO	7
2	Problem Statement	8
2.1	The Filter Trap: Artificial Intelligence in Silos	8
2.2	The Scale Wall: The High Cost of "Pure" Search	8
2.3	Stifled Innovation: The Unbuilt Applications	9
3	Introduction	9
3.1	The Jevons Paradox: Unlocking Dormant Demand	9
3.2	The Structural Crisis	10
3.3	The Cosma Monopoly	10
4	System Architecture	10
4.1	High-Level Overview	10
4.2	Component Interaction	10
5	Cosma Vector Database (CVD)	11
5.1	The Architecture of Efficiency	11
5.2	Proprietary Search Execution	12
5.3	The Hidden Taxes of HNSW (Current Standard)	12
5.4	The Cosma Advantage	13
6	Cosma Omnilingual Embedding Models (COEM)	14
6.1	Model Architecture	14
6.2	Training Methodology	14
6.3	Embedding Space Properties	14
7	Cosma Search Layer (CSL)	14
7.1	End-to-End Search Sequence	15
7.2	Search Features	15
8	Deployment Models	15
9	Security & Compliance	15
9.1	Data Protection	16
9.2	Compliance	16
10	Market Positioning	16
10.1	Competitive Advantages	16
10.2	Target Markets	17
10.3	Competitor Landscape	17
11	Technology Roadmap	19
11.1	Near-Term (6-12 months)	19
11.2	Mid-Term (12-24 months)	19

11.3 Long-Term (24+ months)	19
12 Business Model	19
12.1 Overview	19
12.2 Revenue Streams	19
12.3 Strategic Business Model Advantages	20
12.3.1 1. At-Cost Infrastructure Pass-Through	20
12.3.2 2. Annual Commitments = Predictable Revenue	20
12.3.3 3. Zero-Fee Scaling Flexibility	20
12.3.4 4. Hybrid Pay-As-You-Go + Subscription Model	21
12.3.5 5. Proprietary Ecosystem Lock-in	21
12.3.6 6. Volume Expansion (The Jevons Paradox)	21
12.3.7 7. The "Structure-Free" & "Bridge" Model Moat	22
12.4 COEM Model Pricing Tiers	22
12.5 Customer Segmentation	22
12.6 Unit Economics by Segment	22
12.6.1 Individuals & Hobbyists	23
12.6.2 Small Startups	23
12.6.3 Medium Companies	24
12.6.4 Large Enterprises	24
12.6.5 Technology Giants	25
12.7 Revenue Stream Mix	25
12.8 5-Year Revenue Projection	26
12.9 Key Growth Assumptions	26
12.10 Lean Team Strategy: Capital-Efficient Scaling	26
12.10.1 Headcount Projection	26
12.10.2 How We Stay Lean	27
12.10.3 Benchmark Comparison	27
12.11 Path to Profitability	27
13 Revenue Targets: What We Need to Sell	28
13.1 Year-by-Year Sales Requirements	28
13.1.1 Year 1: \$350K ARR Target (Design Partner Conversions)	28
13.1.2 Year 2: \$5.0M ARR Target	29
13.1.3 Year 3: \$18.7M ARR Target	29
13.2 Embedding Volume Breakdown	30
13.3 Monthly Sales Velocity Requirements	30
13.4 Gross Margin by Revenue Stream	30
13.5 Revenue Sensitivity Analysis	31
14 Intellectual Property	31
14.1 Proprietary Technology	31
14.2 Licensing Strategy	31
15 Risk Factors & Mitigation	31
15.1 Market & Competitive Risks	31
15.1.1 Risk: Competitive Replication	31
15.1.2 Risk: Big Tech Competition	32
15.1.3 Risk: Market Timing	32
15.2 Technical Risks	33

15.2.1 Risk: Scaling Challenges	33
15.2.2 Risk: Model Performance	33
15.3 Execution Risks	33
15.3.1 Risk: Customer Acquisition	33
15.3.2 Risk: Operational Complexity	34
15.4 Regulatory & Compliance Risks	34
15.4.1 Risk: Data Privacy Regulations	34
15.4.2 Risk: AI Liability & Ethics	34
15.5 Financial Risks	34
15.5.1 Risk: Longer Runway Needed	35
15.6 Risk Management Philosophy	35
16 The Ask: \$1.59M Pre-Seed Round	35
16.1 Funding Request	35
16.2 Use of Funds	35
16.3 Key Milestones with This Funding	36
16.3.1 Technical Milestones (Months 1-9)	36
16.3.2 Commercial Milestones (Months 6-18)	36
16.4 Capital Efficiency Strategy	36
16.5 Business Model Moats: Built-In Defensibility	37
16.5.1 1. Self-Funding Growth Engine	37
16.5.2 2. Competitive Immunity from At-Cost Pass-Through	37
16.5.3 3. Expansion Revenue Built Into Architecture	37
16.5.4 4. Technical Moat = Pricing Moat	38
16.6 Post-Series A Vision (18-36 Months)	38
A Competitor Pricing Details	39
A.1 Voyage AI Detailed Pricing	39
A.1.1 Text Embedding Models	39
A.1.2 Multimodal Embeddings	39
A.1.3 Rerankers	39
A.2 Vector Scaling Costs	39
A.3 Search Cost Scaling (1M vs 100M vs 1B Vectors)	39

1 THE TEAM

1.1 Founding Team

Azetta.ai is led by a technical founding team with deep expertise in AI/ML, full-stack engineering, and scalable systems. The team combines academic research credentials with proven entrepreneurial execution.

1.1.1 Taha Bouhsine - Co-Founder & CEO

Background:

- **Google Developer Expert (GDE)** in AI/ML
- **Education:** M.S. in Electrical & Computer Engineering (Rowan University, 2024), M.S. in Data Science, B.Sc in Math & CS
- **Research Fellowship:** Awarded research fellowship at Rowan University based on research initiated in Morocco, enabling continued collaboration with the Federal Aviation Administration (FAA)
- **Research Focus:** Deep Learning, Multimodality, Representation Learning, Continual Learning, and Self-Supervised Learning
- **Publications:** Published research in high-impact factor journals including *IEEE Transactions on Communications (8 Impact Factor)*, with applications to aviation safety, civil engineering, security, and remote sensing

Government-Funded Research Experience:

- **U.S. Federal Agencies:** Worked on research projects funded by the Department of Defense (DoD), Department of Energy (DoE), and Department of Transportation (DoT)
- **International Collaboration:** CNRST Morocco-funded satellite imagery research
- **Multimodal Expertise:** Extensive experience across all modalities (text, vision, audio, satellite imagery, sensor data)

Previous Founding Experience:

- Founded **MLNomads**, a collaborative ML research company. Built an open-source replacement for Jupyter notebooks that scaled to **2,000 unique users**, demonstrating ability to build and scale developer tools
- Founded **SafetyFirst.ai**, a non-profit research lab focused on creating safe foundational models for robotics

Community Leadership & Thought Leadership:

- **Invited Speaker:** As a recognized Google Developer Expert (GDE), regularly invited by international communities and Google to present on AI/ML expertise
- Founded **GDSC UIZ, Tensorflow User Group** in Agadir (90+ events, 2,000+ attendees)
- Organizer for GDG Agadir and GDG Philly

Key Technical Innovations:

- **Kernel Methods Expert:** Pioneered a novel Mercer kernel that eliminates the need for activation functions and layer normalizations in neural networks
- **Primal-Form Kernels:** Developed the first method to use Mercer kernels in their primal

form for neural networks, bypassing traditional dual-form limitations

- **Anchor-Free Contrastive Learning:** Invented a contrastive learning algorithm that removes the need for negative mining, significantly simplifying training dynamics
- **First-Principles Design:** Designs algorithms from mathematical first principles rather than empirical heuristics, ensuring theoretical soundness

Expertise: Taha's research innovations form the mathematical foundation of Cosma's technology stack. His expertise in kernel methods, representation learning, and cross-lingual embeddings, combined with his first-principles approach to algorithm design, positions him uniquely to lead Azetta's technical vision. His ability to translate theoretical breakthroughs into production systems bridges the gap between academic research and commercial impact.

1.1.2 Poppy Yuan - Co-Founder & COO

Background:

- **Education:** M.S. in Operations Research & Supply Chain Management, Case Western Reserve University (Weatherhead School of Management, 2019)
- **Previous Cofounder & COO:** Ekowsystems Inc. (Transportation Platform, 2022-2024)
- **Operations Expertise:** Strategic operations executive with proven track record in scaling early-stage platforms and building customer success infrastructure

Previous Founding Experience - Ekowsystems Inc.:

- Led transportation platform development from concept to beta stage, managing relationships across an **11-person distributed team**, C-level executives, shareholders, and technology partners
- Built end-to-end customer engagement program serving **250+ beta users** with **70% participation rate** through strategic communication and value delivery
- Managed development budget **exceeding \$100K**, evaluating and implementing multiple SaaS solutions
- Implemented comprehensive customer support infrastructure (Twilio, Zendesk) achieving **95% user satisfaction**
- Established systematic project management framework, **reducing feedback cycles by 40%** and ensuring alignment between business requirements and technical execution
- Created market intelligence system to analyze competitor data, user behavior patterns, and industry trends

Prior Experience:

- **Consulting Analyst, Medical Mutual of Ohio (2019-2022):**
 - Analyzed complex datasets using SQL and Tableau, generating insights that drove **\$2.2M annual revenue growth**
 - Managed relationships with 100+ clients, maintaining **100% delivery rate**
 - Implemented automated workflow systems reducing manual work by 20 hours/week and improving reports by 40%
- **Operations Analyst, Cleveland Kitchen (2018-2019):**
 - Implemented lean inventory management system, achieving **50% reduction in re-work** and **20% improvement in forecasting accuracy**
 - Developed standardized digital systems decreasing errors to below 2% and increasing efficiency by 40%

Expertise: Poppy brings proven operational execution from founding a transportation platform to beta stage with strong user engagement metrics. Her expertise in customer success infrastructure, stakeholder management, and data-driven decision-making positions her to scale Azetta's operations from early adopters to enterprise customers. Her experience evaluating and implementing SaaS tools provides deep understanding of customer needs in the B2B SaaS market. Her M.S. in Operations Research ensures strategic, analytical approaches to scaling challenges.

1.1.3 Douglas Seo - Co-Founder & CTO

Background:

- **Education:** B.S. in Electrical Engineering & Computer Science, UC Berkeley (2022)
- **CS 188-189 (Artificial Intelligence & Machine Learning):** Top 5% of class with Professor recognition
- **Relevant Coursework:** Data Structures & Algorithms, Machine Learning, Distributed Systems, Web Development, Cybersecurity

Proven Technical Leadership & Execution:

- **Workflow (2025) - Founding Engineer & CTO**
 - Built AI-powered workflow automation platform, achieving **3,000+ monthly visits**
 - Tech stack: PostgreSQL, React, TypeScript, Supabase, Prisma ORM
 - Architected scalable backend infrastructure supporting real-time data synchronization
 - <https://workflow.us>
- **Popper LLC (December 2022 - Present) - CEO & Founder**
 - Scaled to **3,000+ active users** with **\$400 MRR** and sustainable unit economics
 - Built production React Native application handling **20,000+ daily Firestore operations**
 - Designed and deployed scalable NoSQL database architecture with real-time sync and offline-first capabilities
 - Implemented production-grade REST APIs on Google Cloud with OAuth authentication, analytics pipelines, and messaging infrastructure
 - Led agile engineering team of **8 developers** through code reviews, sprint planning, and technical roadmap execution
 - Achieved profitability with **60+ business clients** as first-time founder
 - <https://popper.social>

Technical Expertise Relevant to Vector Database Development:

- **Database Architecture:** Hands-on experience designing and scaling both SQL (PostgreSQL) and NoSQL (Firestore) systems for production workloads
- **Distributed Systems:** Built real-time synchronization infrastructure handling thousands of concurrent operations
- **Performance Optimization:** Optimized database queries and indexing strategies for sub-100ms latency requirements
- **Cloud Infrastructure:** Production deployment experience on Google Cloud and Supabase with monitoring, logging, and auto-scaling
- **AI/ML Integration:** Academic foundation in machine learning algorithms and practical

experience integrating NLP/speech models into production systems

Value to Cosma: Douglas brings the rare combination of ML/AI academic rigor from UC Berkeley and battle-tested experience building profitable, scalable applications from scratch. His expertise in database architecture, distributed systems, and full-stack engineering ensures Cosma's vector database can evolve from a prototype to production-grade infrastructure capable of serving millions of semantic search queries with enterprise-level reliability.

Bottom Line: This isn't a team learning on investors' dime. All three co-founders have already built products with thousands of real users, secured institutional validation, and demonstrated the ability to execute under resource constraints. Taha's research innovations solve the core technical challenge, Douglas's engineering expertise ensures production-grade execution, and Poppy's operational excellence guarantees customer success at scale. The combination of deep technical moat, proven execution, and institutional trust makes this an exceptionally strong founding team for a deep-tech pre-seed venture.

2 PROBLEM STATEMENT

The AI infrastructure market is approaching a critical **economic cliff**. As enterprises transition from building prototypes to deploying global-scale applications, they are discovering that the current generation of vector databases is mathematically incapable of scaling without destroying unit economics.

The Core Dilemma: Scale vs. Scope

Today's architects are forced to choose between two fatal options:

1. **Limit the Scope (The Filter Trap):** Restrict search to specific languages or metadata tags to keep costs low, but sacrifice global intelligence.
2. **Pay the Scale Tax (The RAM Cliff):** Attempt "pure" search over the entire dataset, requiring massive RAM usage that makes unit economics impossible.

2.1 The Filter Trap: Artificial Intelligence in Silos

To survive existing infrastructure costs, 95% of enterprise deployments rely on **language-specific filters**. They artificially partition data—isolating "English Customer Support" from "Spanish Technical Docs"—to reduce the search space.

This is a functional regression. It re-imposes the very silos that vector search was promised to destroy. A query in Hindi cannot leverage knowledge from an English manual because the system is explicitly told not to look there. The cost is lowered, but the **intelligence is capped**.

2.2 The Scale Wall: The High Cost of "Pure" Search

"Pure" vector search—scanning the high-dimensional manifold to find the absolute best match regardless of partition—is the holy grail of retrieval. However, it requires an infrastructure that does not exist today.

Existing solutions (HNSW, IVFPQ) suffer from a crippling mathematical flaw: **Linear Cost Scaling** ($O(N)$).

- **The RAM Cliff:** To deliver sub-second latency, indices must be held in RAM. As data grows to billions of vectors, memory costs scale non-linearly, forcing companies to buy expensive, high-memory GPU instances just to store the index.
- **The Accuracy Tax:** To avoid this cost, systems use aggressive compression (quantization), which degrades result quality. You pay less, but your AI gets dumber.

2.3 Stifled Innovation: The Unbuilt Applications

This infrastructure gap is not just an inconvenience; it effectively **bans** the creation of high-value applications. A class of transformative products remains theoretically possible but economically unbuildable:

- **Global Information Engines:** Systems that instantly synthesize insights from a corporation's *total* knowledge base, across 100+ languages, without manual routing.
- **Planetary-Scale Biometrics:** Face recognition authentication that functions seamlessly across billions of identities, vital for global security and payments.
- **Universal Genomic Matching:** Searching billions of genetic sequences to identify rare disease markers or patient matches without pre-filtering by population or study type.
- **Real-Time Content Integrity:** Instantly checking every frame of user-uploaded video against a global media registry for copyright or safety violations, impossible with today's batch-processing lag.
- **Unbounded Pattern Discovery:** Fraud detection systems that find subtle correlations across petabytes of unstructured logs without needing pre-defined "suspect" filters.

The Opportunity

The market does not need another incremental optimization. It needs a **first-principles infrastructure overhaul**. We provide the foundation that eliminates this "Filter vs. Scale" trade-off, enabling pure, global vector search to become the profitable standard for the industry.

3 INTRODUCTION

We stand at the precipice of a new era in human knowledge: the **Omnilingual Age**. The artificial barriers that have divided information for millennia—language, script, and geography—are collapsing. In this new world, a query in English must instantly unlock insights in Mandarin, Arabic, or Hindi.

This is not a feature request; it is an inevitability. But the infrastructure to support it does not exist.

3.1 The Jevons Paradox: Unlocking Dormant Demand

The **Jevons Paradox** states that increasing the efficiency of a resource leads to an explosion in its consumption. Today, semantic search is the world's most valuable resource, yet it remains

artificially scarce. It is capped by a broken legacy stack that imposes a "Multilingual Tax" on every query and a "Vector Tax" on every byte of data.

Cosma is the catalyst designed to shatter these caps. We are not merely building a faster database; we are engineering the efficiency breakthrough that will trigger the Jevons Paradox for global intelligence.

3.2 The Structural Crisis

The current search paradigm is facing an existential crisis. As organizations race to index billions of vectors across hundreds of languages, they are hitting a mathematical wall:

- **The Fragmentation Trap:** Maintaining 100+ separate indices for different languages is operationally impossible at scale.
- **The RAM Cliff:** Relying on memory-heavy graph indices (HNSW) makes high-scale search economically suicidal.
- **The Accuracy Ceiling:** Approximate methods sacrifice the very precision that enterprise AI demands.

3.3 The Cosma Monopoly

Cosma is the world's first vertically integrated **Omnilingual Semantic Search Platform**. We have rebuilt the stack from first principles to create a defensible monopoly on the future of search:

1. **Unified Intelligence (COEM):** One vector space for all languages. No translation, no silos.
2. **Infinite Scale (CVD):** A proprietary disk-native engine that turns storage into search, eliminating the RAM bottleneck.
3. **Frictionless Access (CSL):** A production-ready layer that democratizes this power for every developer.

We are not competing for a slice of the existing search market. We are unlocking a new, exponentially larger market—one where the cost of intelligence falls to zero, and the volume of consumption scales to infinity.

4 SYSTEM ARCHITECTURE

4.1 High-Level Overview

Figure 1 illustrates Cosma's four-layer architecture designed for enterprise deployment.

4.2 Component Interaction

Figure 2 shows detailed relationships between core components.




Figure 1: Cosma System Architecture - Four-Layer Design

5 COSMA VECTOR DATABASE (CVD)

We utilize a proprietary storage engine that eliminates graph-maintenance overhead.

5.1 The Architecture of Efficiency

Traditional vector databases rely on graph-based indexing (HNSW) or cluster-based approximations (IVF) to overcome the scaling limit. These legacy architectures introduce unavoidable hardware and operational taxes.

What We Eliminated

- **No Graph Construction Overhead:** Competitors burn 30-50% of CPU cycles just “wiring” indices. CVD eliminates this step, enabling instant availability.
- **No “RAM Cliff”:** Graph indices must reside in memory to be performant, forcing customers to rent expensive high-RAM instances (e.g., AWS r6g.24xlarge). CVD is archi-




Figure 2: Component Interaction Diagram - Data Flow Architecture

tested for Disk-Native Efficiency, running on standard commodity storage without latency degradation.

- **No Re-Indexing Lag:** CVD removes the need for background “build” or “warm-up” phases. Data is searchable the millisecond it is ingested.

5.2 Proprietary Search Execution

CVD employs a trade-secret Direct-Storage Traversal algorithm. By decoupling semantic precision from RAM-heavy graph structures, we achieve:

- **Predictable Latency:** Removing the probabilistic variance of graph “hops.”
- **Linear Cost Scaling:** Costs grow with storage (cheap), not memory (expensive).

5.3 The Hidden Taxes of HNSW (Current Standard)

While HNSW (used by Pinecone, Weaviate) offers fast read speeds, it forces a “Trilemma” of compromises that Cosma eliminates:




Figure 3: CVD Query Execution Pipeline - Deterministic Search Flow

- **The Write Tax:** HNSW requires calculating distances to multiple neighbors and rewriting graph edges for every insertion. *Result:* Write throughput is capped, and CPUs spike during ingestion.
- **The Memory Tax:** The HNSW graph structure requires storing millions of edge pointers in RAM. *Result:* 1 Billion vectors require $\sim 1.5\text{TB} - \sim 3\text{TB}$ of RAM, costing $> \$15\text{k/month}$ in raw infrastructure just to hold the index.
- **The Mutability Tax:** Deleting or updating vectors in a graph is computationally expensive and often degrades the graph's quality over time, forcing a complete "Re-Index" (downtime/CPU spike).

5.4 The Cosma Advantage

Cosma eliminates the graph entirely. Our proprietary indexing method is stateless regarding neighbor relationships at the storage layer. This means:

- **O(1) Memory Overhead:** We do not store edge pointers.
- **Instant Mutability:** Updates and deletes are standard database operations, not graph repairs.

“Unlike graph-based approximations where query pathing changes based on index fragmentation (leading to jitter), Cosma’s search path is architecturally fixed. This delivers consistent latency profiles regardless of data scale.”

6 COSMA OMNILINGUAL EMBEDDING MODELS (COEM)

COEM models solve the fundamental challenge of cross-lingual semantic representation by mapping 100+ languages into a unified vector space where semantically similar content clusters regardless of language.

6.1 Model Architecture




Figure 4: COEM Model Architecture - Unified Cross-Lingual Latent Space

6.2 Training Methodology

COEM models achieve cross-lingual alignment through:

1. **Parallel Corpus Training:** Millions of aligned text pairs across 100+ languages
2. **Contrastive Learning:** Pushes semantically similar content together, dissimilar apart
3. **Manifold Regularization:** Ensures smooth geometric structure in latent space
4. **Domain-Specific Augmentation:** Fine-tuning for search-optimized representations

6.3 Embedding Space Properties

7 COSMA SEARCH LAYER (CSL)

CSL provides the production-ready search API that coordinates embedding generation, vector search, and result ranking.




Figure 5: Unified Embedding Space - Semantic Clustering Across 100+ Languages




Figure 6: CSL Search Sequence - End-to-End Query Processing Flow

7.1 End-to-End Search Sequence

7.2 Search Features

- **Automatic Embedding:** Text automatically converted to vectors
- **Metadata Filtering:** Structured filters (date, category, author)
- **Semantic Re-ranking:** Distance + coherence + language priors
- **Result Normalization:** Consistent scoring across languages

8 DEPLOYMENT MODELS

Cosma supports three deployment tiers designed for different organizational needs.

9 SECURITY & COMPLIANCE




Figure 7: Cosma Deployment Tiers - Progressive Service Levels

9.1 Data Protection

- **Encryption:** All data encrypted in transit (TLS 1.3) and at rest (AES-256)
- **Isolation:** Namespace-level separation, dedicated compute for premium
- **Privacy:** Customer data never used for model training
- **Logging:** Anonymized logs, configurable retention

9.2 Compliance

Cosma meets enterprise compliance requirements:

- SOC 2 Type II (in progress)
- GDPR compliant
- HIPAA-ready (enterprise tier)
- ISO 27001 certified data centers

10 MARKET POSITIONING

10.1 Competitive Advantages

1. **True Omnilingual Search:** Single unified index, not language-specific routing

2. **Deterministic Performance:** No ANN approximations = predictable latency
3. **Enterprise Security:** SOC 2 Type II compliance and dedicated isolated environments
4. **Cost Efficiency:** CPU-optimized, 60% lower infrastructure costs vs. GPU-dependent solutions

10.2 Target Markets

- **Global Enterprises:** Companies with multilingual content (legal, customer support, documentation)
- **E-commerce:** Cross-border product search and recommendations
- **Healthcare:** Medical literature search across languages
- **Government:** Intelligence, diplomatic communications, public records
- **Legal:** Contract search, due diligence, compliance

10.3 Competitor Landscape

The semantic search market includes several key players. Cosma differentiates itself through its unified omnilingual architecture and deterministic performance.

- **Voyage AI:**
 - *Business Model:* Proprietary SaaS API provider specializing in embeddings and reranking.
 - *Pricing Model:* Consumption-based. Generous free tier (first 200M tokens free). Paid usage ranges from \$0.02 to \$0.18 per 1 million tokens depending on the model. Multimodal and Reranking models have separate pricing structures.
 - *Focus:* High-quality embedding models and rerankers, particularly for RAG applications.
 - *Target Customers:* Developers and AI-native companies building RAG pipelines.
 - *Differentiation:* While Voyage excels in model quality, it lacks a native vector DB.
- **Weaviate:**
 - *Business Model:* Open-core model (Open Source software + Managed Cloud services).
 - *Pricing Model:* Hybrid: Free open-source; Serverless (pay-as-you-go); Enterprise.
 - *Focus:* AI-native open-source vector database with a modular architecture.
 - *Target Customers:* Broad range from startups to enterprises requiring flexibility and modular integrations.
 - *Differentiation:* Weaviate relies on HNSW (ANN) for search, which introduces probabilistic recall. Cosma's CVD uses a deterministic algorithm, ensuring 100% recall accuracy.
- **Pinecone:**
 - *Business Model:* Closed-source fully managed SaaS (Database-as-a-Service).
 - *Pricing Model:* Usage-based (Serverless read/write units) and Provisioned (hourly rate per pod/instance).
 - *Focus:* Fully managed, cloud-native vector database known for ease of use.

- *Target Customers:* Developers and enterprises seeking a “hands-off” managed infrastructure.
 - *Differentiation:* Pinecone is a general-purpose vector DB. Cosma offers a vertically integrated solution (Model + DB) specifically optimized for cross-lingual semantic search.
- **pgvector:**
 - *Business Model:* Free and Open Source (PostgreSQL extension).
 - *Pricing Model:* Free software; costs are incurred only for the underlying PostgreSQL infrastructure (self-hosted or managed).
 - *Focus:* Open-source vector similarity search extension for PostgreSQL.
 - *Target Customers:* Existing PostgreSQL users and smaller-scale applications.
 - *Differentiation:* pgvector is excellent for simplicity but struggles with performance and latency at enterprise scale (100M+ vectors). Cosma’s CVD is purpose-built for high-scale vector workloads.
 - **Milvus / Zilliz:**
 - *Business Model:* Open-core (Open Source Milvus + Zilliz Cloud managed service).
 - *Pricing Model:* Zilliz Cloud: Consumption-based (compute units + storage). Free tier available.
 - *Focus:* Scalability-first vector database designed for billion-scale datasets.
 - *Target Customers:* Large-scale AI applications, research institutions, enterprises with massive vector workloads.
 - *Differentiation:* Milvus relies on **memory-heavy HNSW graphs** (requiring 2-3x RAM vs raw data) or complex disk-based graphs (DiskANN/Vamana) to function. Azetta’s **“Structure-Free” architecture** eliminates the graph entirely, delivering comparable scale at **5% of the infrastructure cost** by natively sorting data on disk without RAM-intensive indexes.
 - **Elasticsearch (with vector search):**
 - *Business Model:* Open-core (Elasticsearch open source + Elastic Cloud).
 - *Pricing Model:* Elastic Cloud: Tiered subscriptions (Standard, Gold, Platinum, Enterprise).
 - *Focus:* The incumbent full-text search platform, now adding vector search capabilities.
 - *Target Customers:* Enterprises with existing Elasticsearch deployments; observability and logging use cases.
 - *Differentiation:* Elasticsearch is not **native vector**—it bolted on vector search to a document store. Performance degrades at high vector density. Cosma is vector-native from the ground up, optimized for semantic search workloads. Additionally, enterprises using Elastic often face **licensing complexity** (Elastic License 2.0 vs OpenSearch fork).

Margin Note: Our at-cost infrastructure pass-through creates competitive immunity on price. However, our real margin comes from **COEM model usage fees (80-90% gross margin)**, not infrastructure. This is the inverse of competitors who mark up infrastructure but commoditize embeddings.

11 TECHNOLOGY ROADMAP

11.1 Near-Term (6-12 months)

- Public SDK release (Python, JavaScript, Go)
- Expanded language support (120+ languages)
- Index compression (8-bit quantization)
- Real-time indexing pipeline

11.2 Mid-Term (12-24 months)

- Distributed sharding for 100B+ vectors
- Streaming search capabilities
- Multi-modal embeddings (text + images)
- Edge deployment (low-power devices)

11.3 Long-Term (24+ months)

- Open adapter ecosystem for domain-specific models
- Plugin architecture for custom distance functions
- Federated search across multiple CVD instances
- AutoML for embedding model customization

12 BUSINESS MODEL

12.1 Overview

Azetta.ai delivers **Cosma**, an enterprise-grade omnilingual semantic search platform operating on a tiered subscription model with usage-based pricing. We target five distinct customer segments with differentiated value propositions and cost structures, generating revenue through three primary streams.

12.2 Revenue Streams

1. **Model Usage Fees:** Consumption-based pricing for COEM embeddings
 - Dimension-based tiered pricing (see Section 12.4)
 - Bundled in subscription tiers or pay-as-you-go overage
2. **Managed Infrastructure:** At-cost pass-through pricing
 - *Cloud Cost:* Customer pays raw cloud provider rate (e.g., AWS/GCP instance cost) with no markup.
 - *Billing:* Annual commitment preferred.
 - *Flexibility:* Zero-fee instance scaling (upgrade/downgrade) enabled by decoupled storage architecture.
3. **Professional Services:** Integration, training, custom models
 - Typical engagements: \$50K-\$250K
 - Expected revenue mix: 10-15% (Years 1-2), declining to 5% (Year 5)

- Services include: Custom model training, integration support, enterprise onboarding

12.3 Strategic Business Model Advantages

Azetta.ai's pricing model is architected to create **sustainable competitive moats** while maximizing customer lifetime value:

12.3.1 1. At-Cost Infrastructure Pass-Through

What it means: Customers pay the raw cloud provider cost (AWS/GCP) for compute and storage with **zero markup**.

Why this wins:

- **Radical Transparency:** No hidden margins or surprise infrastructure bills—customers see exactly what they're paying for
- **Competitive Immunity:** Competitors can't undercut us on infrastructure since we charge zero margin
- **Trust Builder:** Enterprise buyers love knowing we're not profiting from their infrastructure spend—builds long-term partnerships
- **Revenue Focus:** We make money on value-add (COEM models, search intelligence), not commodity infrastructure

Customer impact: A company spending \$5,000/month on infrastructure pays exactly \$5,000. With traditional vendors (20-40% markup), the same workload costs \$6,000-\$7,000/month.

12.3.2 2. Annual Commitments = Predictable Revenue

What it means: Infrastructure costs are billed **annually in advance**, not month-to-month.

Why this wins:

- **12-Month Cash Flow:** Every customer provides upfront capital, reducing working capital needs
- **Lower Churn Risk:** Annual contracts create switching friction—customers don't casually cancel mid-year
- **Financial Predictability:** ARR becomes highly predictable, making revenue forecasting accurate for investors
- **Customer Savings:** We pass along AWS annual commitment discounts (10-15% savings vs. on-demand)

Investor impact: By Year 3, 70%+ of revenue is under annual contracts, providing **\$19.7M+ in prepaid cash** to fund operations without relying on external capital.

12.3.3 3. Zero-Fee Scaling Flexibility

What it means: Customers can scale infrastructure up or down **without change fees or penalties**, enabled by our decoupled storage architecture.

Why this wins:

- **Eliminates Upgrade Friction:** Customers don't fear scaling up (no "gotcha" fees), accelerating expansion revenue
- **Technical Moat:** Decoupled architecture is hard to replicate—competitors use fixed instances
- **Expansion Revenue Driver:** Easy scaling drives 20-30% annual upsells through natural usage growth
- **Risk Mitigation:** Customers can downgrade during slowdowns without penalty, reducing churn during economic downturns

Economic impact: Net Revenue Retention (NRR) of 110-130% driven by frictionless expansion—every \$1 of new ARR becomes \$1.10-\$1.30 the following year.

12.3.4 4. Hybrid Pay-As-You-Go + Subscription Model

What it means: Customers buy bundled tiers (predictable base) + pay overage fees for usage spikes (consumption alignment).

Why this wins:

- **Land-and-Expand:** Start with small subscription, grow into higher tiers as usage increases organically
- **Revenue Upside:** Overage fees capture unexpected growth (20-30% of customers exceed base tier monthly)
- **Budget Certainty:** Base subscription is fixed; overage only applies when customers scale, aligning costs with their revenue growth
- **Self-Service Friendly:** Usage-based pricing removes sales friction—customers self-select appropriate tiers

Revenue mix: 60% subscription (predictable) + 30% infrastructure pass-through (annual) + 10% overage/services (expansion) = highly diversified revenue base.

12.3.5 5. Proprietary Ecosystem Lock-in

What it means: The Cosma Vector Database (CVD) is tightly coupled with our model architecture.

Why this wins:

- **Technical Dependency:** CVD's "zero-build" efficiency relies on specific vector properties produced only by Cosma models or our adaptation API.
- **No "Bring Your Own Vector" Leakage:** Customers cannot simply generate vectors with OpenAI and query CVD directly. They *must* use our adaptation layer, ensuring we capture revenue even if they use third-party models.
- **Sticky Ecosystem:** The database and model reinforce each other—using one makes the other more valuable.

12.3.6 6. Volume Expansion (The Jevons Paradox)

What it means: By reducing the unit cost of vector search by 10x, we unlock massive-scale use cases that were previously cost-prohibitive.

Why this wins:

- **New Workloads:** Instead of just indexing “hot” data (recent 10%), customers can now afford to index *everything* (archival data, logs, full history).
- **Competitive Defensibility:** Pass-through pricing + decoupled architecture = nearly impossible for competitors to match without rebuilding their entire platform

12.3.7 7. The “Structure-Free” & “Bridge” Model Moat

What it means: Current vector databases require massive RAM because they must store complex graphs (HNSW) to understand similarity. Azetta has developed the world’s first **“Structure-Free” Embedding Model**. Our proprietary model embeds data directly into a format that standard storage engines can sort natively.

Why this wins:

- **5% Cost Basis:** This eliminates the need for the graph entirely, allowing us to run at 5% of the cost of Pinecone with mathematical guarantees on retrieval speed.
- **The “Bridge” Model Monopoly:** You cannot use standard Postgres for this because standard embeddings (from OpenAI, Cohere, or Voyage) are mathematically incompatible with disk-based indexing. If you try to sort OpenAI vectors on a disk, you get garbage results.
- **Zero-RAM Search:** Azetta is not just a database wrapper. We are the only provider of the necessary “bridge” model—the COEM—that translates human language into this high-efficiency format. Customers pay us because we own the only model in the world that makes “zero-RAM” search possible.

12.4 COEM Model Pricing Tiers

COEM models are available in multiple dimensions, with pricing reflecting computational cost. Subscription tiers bundle prepaid COEM tokens with managed infrastructure and unlimited search queries.

Model	Dims	\$/1M Tokens	Rel. Cost	Use Case
COEM-Tiny	128	\$0.02	1x	High-volume, basic similarity
COEM-Small	256	\$0.04	2x	Standard applications
COEM-Base	512	\$0.08	4x	Production workloads
COEM-Large	768	\$0.12	6x	High-accuracy search
COEM-XL	1024	\$0.18	9x	Premium semantic understanding
COEM-XXL	1536	\$0.30	15x	Maximum precision

Table 1: COEM Pricing by Dimension (100+ Languages, Unlimited Searches)

Pricing factors: Higher dimensions = more GPU memory, larger index storage, lower throughput per GPU, but better semantic precision.

Example: The “Professional” tier (\$999/month) includes 5M embeddings (40B tokens using COEM-Base at $\$0.08/1M = \3.20 in compute) plus hosted infrastructure and unlimited searches.

12.5 Customer Segmentation

12.6 Unit Economics by Segment

Segment	Typical Use Case	Infrastructure Tier	Monthly Embeddings	Price Range
Individuals/ Hobbyists	Personal projects, learning, prototyping	Shared CPU	10K-100K	Free/\$9-29
Small Startups	MVP development, early apps with <10K users	1-2 vCPU instance	100K-1M	\$99-299
Medium Companies	Production apps, 10K-100K users	4-8 vCPU instance	1M-10M	\$999-2,999
Large Enterprises	Multiple products, 100K-1M users	Dedicated cluster (16-32 vCPU)	10M-100M	\$9,999-29,999
Tech Giants	Platform-scale, millions of users	Multi-region clusters	100M+	Custom (6 figures)

Table 2: Customer Segment Characteristics (Unlimited Searches on Hosted Infrastructure)

12.6.1 Individuals & Hobbyists

Revenue Model:

- Free Tier: Shared CPU + 10,000 embeddings/month (unlimited searches, \$0 revenue)
- Starter: \$9/month - Shared CPU + 50,000 embeddings (unlimited searches)
- Pro: \$29/month - 1 vCPU + 200,000 embeddings (unlimited searches)
- **Included Model:** COEM-Small (256 dims, \$0.04/1M tokens)
- Overage: \$0.10 per 10K embeddings beyond quota (COEM-Small)

Cost Analysis (per paying user):

- Infrastructure (shared compute): \$2-4/month
- Model inference (GPU, batched): \$1-3/month
- Storage & bandwidth: \$0.50-2/month
- **Total COGS:** \$3.50-9/month
- **Gross margin:** 0-70% (Free tier subsidized by paid users)

Conversion assumptions: 5-10% of free users convert to paid tiers.

12.6.2 Small Startups

Revenue Model:

- Growth: \$99/month - 1 vCPU instance + 500K embeddings (unlimited searches)
- Business: \$299/month - 2 vCPU instance + 2M embeddings (unlimited searches)
- **Included Model:** COEM-Small (256 dims, \$0.04/1M tokens)
- Overage: \$0.10 per 10K embeddings beyond quota (COEM-Small)

Cost Analysis (per customer):

- Infrastructure (1-2 vCPU): \$20-40/month
- Model inference (GPU): \$10-25/month
- Storage & bandwidth: \$5-10/month

- Support overhead: \$5-10/month (email support)
- **Total COGS:** \$40-85/month
- **Gross margin:** 60-75%

Revenue projections:

- Average deal size: \$150-250/month (including overage)
- CAC: \$500-1,000 (self-service, marketing-driven)
- Payback period: 4-6 months
- Annual LTV: \$1,800-3,000 (assuming 12-18 month retention)

12.6.3 Medium Companies

Revenue Model:

- Professional: \$999/month - 4 vCPU instance + 5M embeddings (unlimited searches)
- Premium: \$2,999/month - 8 vCPU instance + 20M embeddings (unlimited searches)
- **Included Model:** COEM-Base (512 dims, \$0.08/1M tokens)
- Overage: \$0.12 per 10K embeddings (COEM-Base)
- Includes: 99.9% SLA, dedicated support, advanced features

Cost Analysis (per customer):

- Infrastructure (4-8 vCPU): \$200-400/month
- Model inference (GPU): \$100-250/month
- Storage & bandwidth: \$50-100/month
- Support & account management: \$100-150/month
- **Total COGS:** \$450-900/month
- **Gross margin:** 65-75%

Revenue projections:

- Average deal size: \$1,500-3,500/month
- CAC: \$5,000-10,000 (direct sales)
- Payback period: 4-6 months
- Annual LTV: \$30,000-50,000

12.6.4 Large Enterprises

Revenue Model:

- Enterprise: \$9,999/month - 16 vCPU cluster + 50M embeddings (unlimited searches)
- Enterprise Plus: \$29,999/mo - 32 vCPU + 200M embeddings (unlimited searches)
- **Included Model:** COEM-Large (768 dims, \$0.12/1M tokens)
- Overage: \$0.15 per 10K embeddings (COEM-Large)
- Includes: Single-tenant, VPC peering, premium support, compliance

Cost Analysis (per customer):

- Dedicated infrastructure (16-32 vCPU): \$2,000-5,000/month
- GPU compute cluster: \$1,500-3,500/month
- Storage & bandwidth: \$500-1,000/month

- Premium support & CSM: \$1,000-2,000/month
- **Total COGS:** \$5,000-11,500/month
- **Gross margin:** 50-70%

Revenue projections:

- Average deal size: \$15,000-35,000/month
- CAC: \$30,000-60,000 (enterprise sales)
- Payback period: 6-12 months
- Annual LTV: \$200,000-450,000
- Contract length: 1-3 years

12.6.5 Technology Giants

Revenue Model:

- Custom pricing: \$100,000-500,000+/month
- **Included Model:** COEM-XL or COEM-XXL (1024-1536 dims, negotiated)
- Overage: Custom volume-based pricing
- Volume discounts, custom SLAs, white-glove service
- Platform integration, co-marketing opportunities

Cost Analysis (per customer):

- Massive dedicated infrastructure: \$30,000-100,000/month
- GPU fleets: \$20,000-80,000/month
- Storage & global CDN: \$5,000-20,000/month
- Dedicated engineering & support: \$10,000-25,000/month
- **Total COGS:** \$65,000-225,000/month
- **Gross margin:** 35-55%

Revenue projections:

- Average deal size: \$150,000-400,000/month
- CAC: \$100,000-250,000 (strategic sales)
- Payback period: 12-18 months
- Annual LTV: \$2M-5M+
- Contract length: 2-5 years

12.7 Revenue Stream Mix

Our revenue model evolves as the business scales, with infrastructure and services revenue growing alongside core model usage:

Revenue Stream	Years 1-2	Years 3-5
Model Usage Fees	79%	55%
Managed Infrastructure	8%	35%
Professional Services	13%	10%

Table 3: Projected Revenue Mix by Stream

Infrastructure revenue grows as customers scale to larger deployment tiers, while professional services remain stable at 10% as improved documentation and self-service tooling reduce

dependency.

12.8 5-Year Revenue Projection

Year 1 Reality Check: Enterprise sales cycles are 9-15 months. Year 1 revenue comes from **design partner conversions** and SMB self-service—not enterprise deals. Enterprise ARR materializes in Year 2 once the sales pipeline matures.

Segment	Year 1	Year 2	Year 3	Year 4	Year 5
Individuals (paid)	\$20K	\$100K	\$300K	\$600K	\$1.0M
Small Startups	\$180K	\$800K	\$2.0M	\$4.0M	\$7.0M
Medium Companies	\$120K	\$1.0M	\$5.0M	\$10M	\$18M
Large Enterprises	\$30K	\$1.3M	\$7.2M	\$18M	\$35M
Tech Giants	\$0	\$600K	\$4.2M	\$12M	\$24M
Total ARR	\$350K	\$3.8M	\$18.7M	\$44.6M	\$85M

Table 4: Projected Annual Recurring Revenue by Segment (Conservative)

12.9 Key Growth Assumptions

1. Customer acquisition ramps progressively: 60 → 400 → 1,500 → 3,500 → 6,000 total customers
2. Year 1 is primarily design partner conversions; enterprise sales motion begins Month 10
3. Average revenue per customer grows from \$486/month (Y1) to \$1,180/month (Y5)
4. Gross margin improves from 55% (Y1) to 70% (Y5) through scale efficiencies
5. **Churn rates:** 5%/mo (hobbyists), 3%/mo (SMB), 1%/mo (enterprise), 0.5%/mo (tech giants)
6. **Net Revenue Retention:** 110-130% driven by usage growth and tier upgrades
7. Infrastructure costs decrease 15-20% annually through optimization and better pricing
8. Enterprise sales cycle: 9-15 months (Year 1 enterprise revenue is pilot conversions only)

12.10 Lean Team Strategy: Capital-Efficient Scaling

Unlike traditional SaaS companies that aggressively hire large sales teams, Azetta.ai is designed for **extreme operational leverage** through automation, product-led growth, and technical excellence.

12.10.1 Headcount Projection

Function	Year 1	Year 2	Year 3	Year 4	Year 5
Engineering (Product & Infra)	3	5	8	12	15
Sales & Marketing	1	2	4	6	8
Customer Success & Support	0.5	2	4	6	8
G&A (Finance, Legal, Ops)	0.5	1	2	3	6
Total Headcount	5	10	18	27	37
Revenue per Employee	\$70K	\$500K	\$1.04M	\$1.65M	\$2.3M

Table 5: Projected Headcount & Revenue per Employee

Year 2 Support Scaling: We add a **Technical Support Engineer** in Year 2 to handle the jump from 60 to 400 customers. This brings customer:support ratio to 1:200 (vs. industry standard 1:30)—aggressive but achievable with strong documentation, self-service tooling, and technical buyers who prefer docs over phone support.

12.10.2 How We Stay Lean

1. **Product-Led Growth (PLG):** Self-service onboarding and freemium tiers eliminate need for large SDR/BDR teams. Target: 70%+ customers acquired through product, not outbound sales
2. **Automated Customer Success:** Built-in usage analytics, proactive alerts, and comprehensive documentation reduce support tickets by 60%. One CSM can manage 100+ accounts vs. industry standard of 30-40
3. **Engineering Leverage:** Investing in infrastructure automation (auto-scaling, monitoring, deployment) means 15 engineers in Year 5 can support \$122M ARR workloads that competitors need 50+ engineers to manage
4. **Founder-Led Sales (Early Stage):** Founders handle enterprise deals through Year 2, deferring expensive sales hires until repeatable playbook is proven
5. **Outsourced Non-Core Functions:** Accounting, legal, HR, and recruiting handled via fractional executives and external vendors rather than full-time hires

12.10.3 Benchmark Comparison

Metric	Azetta.ai	Typical SaaS	Advantage
Revenue per Employee (Year 5)	\$3.5M	\$150K-250K	10-20x
Customer Support Ratio	1:100+	1:30-40	3x leverage
Sales Team % of Headcount	23%	40-50%	50% reduction
Eng % of Headcount	43%	25-30%	Product focus

Table 6: Lean Operations vs. Industry Benchmarks

Capital Efficiency Impact: By Year 5, Azetta.ai achieves \$2.3M revenue per employee—**8-15x higher than typical SaaS companies**. This operational leverage positions us to reach profitability by Year 4 with less dilution.

12.11 Path to Profitability

Metric	Year 1	Year 2	Year 3	Year 4	Year 5
Total Revenue	\$350K	\$3.8M	\$18.7M	\$44.6M	\$85M
COGS (45-30%)	-\$158K	-\$550K	-\$5.99M	-\$13.4M	-\$25.5M
Gross Profit	\$192K	\$3.25M	\$12.71M	\$31.2M	\$59.5M
Gross Margin	55%	85%	68%	70%	70%
Sales & Marketing (as % revenue)	-\$245K (70%)	-\$2.5M (65%)	-\$6.55M (35%)	-\$11.15M (25%)	-\$17M (20%)
R&D (as % revenue)	-\$140K (40%)	-\$1.75M (46%)	-\$5.61M (30%)	-\$11.15M (25%)	-\$17M (20%)
G&A (as % revenue)	-\$53K (15%)	-\$600K (16%)	-\$2.24M (12%)	-\$5.35M (12%)	-\$10.2M (12%)
EBITDA	-\$246K	-\$1.6M	-\$1.69M	+\$3.55M	+\$15.3M
EBITDA Margin	-70%	-42%	-9%	+8%	+18%

Table 7: Projected Income Statement & Path to Profitability

Break-even: EBITDA-positive by Year 4 at \$44.6M ARR. Company prioritizes aggressive market share capture in Years 1-3 while steadily improving operational efficiency. By Year 5, achieves best-in-class SaaS metrics with 70% gross margins and 18% EBITDA margins, demonstrating sustainable unit economics at scale.

Key efficiency drivers:

- **S&M optimization:** Product-led growth (PLG) motion reduces CAC by 60% from Year 1 to Year 5; self-service adoption drives 70%+ of new customers by Year 5
- **R&D leverage:** Core platform matures by Year 3; later-stage R&D focuses on incremental features and vertical expansion rather than foundational rebuilds
- **Infrastructure efficiency:** Economies of scale reduce per-customer compute costs 40% through better GPU utilization and AWS volume discounts

Note: This trajectory balances growth and profitability. More aggressive growth scenario (maintaining 30% S&M spend through Year 5) could achieve \$200M+ ARR but delay profitability to Year 6.

13 REVENUE TARGETS: WHAT WE NEED TO SELL

This section provides **exact unit economics** to achieve our financial projections. Revenue is driven by three streams: COEM model usage (embeddings), managed infrastructure (VMs), and professional services.

Important: Year 1 revenue is primarily from **design partner conversions**, not enterprise sales. Enterprise sales cycles are 9-15 months, so Year 1 enterprise ARR reflects pilot conversions only.

13.1 Year-by-Year Sales Requirements

13.1.1 Year 1: \$350K ARR Target (Design Partner Conversions)

Revenue Stream	Mix	Target	Unit Price	Units Required
Model Usage Fees (60% = \$210K/year)				
COEM-Small (256d)	50%	\$105K	\$0.04/1M tokens	2.63B tokens
COEM-Base (512d)	35%	\$73.5K	\$0.08/1M tokens	919M tokens
COEM-Large (768d)	15%	\$31.5K	\$0.12/1M tokens	263M tokens
Total Tokens				3.8B/year
Monthly Average				317M tokens/month
Managed Infrastructure (30% = \$105K/year)				
1-2 vCPU Instances	60%	\$63K	\$30/mo avg	175 instance-months
4-8 vCPU Instances	35%	\$37K	\$300/mo avg	123 instance-months
16+ vCPU Clusters	5%	\$5.25K	\$3,500/mo avg	1.5 cluster-months
Avg Monthly Infrastructure				\$8.75K/month
Professional Services (10% = \$35K/year)				
Integration Support	100%	\$35K	\$10K avg	3-4 engagements

Table 8: Year 1 Sales Targets by Revenue Stream (Conservative)

Year 1 Customer Mix (60 customers total):

- **30 Hobbyist/Starter** @ \$55/mo avg = \$20K ARR (5.7%)
- **20 Small Startups** @ \$750/mo avg = \$180K ARR (51.4%)
- **8 Medium Companies** @ \$1,250/mo avg = \$120K ARR (34.3%)
- **2 Large Enterprises (pilots)** @ \$1,250/mo avg = \$30K ARR (8.6%)
- **0 Tech Giants** = \$0 (0%)

Key Insight: Year 1 enterprise revenue (\$30K) comes from pilot conversions at discounted rates, not full enterprise contracts. This is realistic given 9-15 month sales cycles.

13.1.2 Year 2: \$5.0M ARR Target

Revenue Stream	Mix	Target	Unit Price	Units Required
Model Usage Fees (60% = \$3.0M/year)				
COEM-Small (256d)	30%	\$900K	\$0.04/1M tokens	22.5B tokens
COEM-Base (512d)	35%	\$1.05M	\$0.08/1M tokens	13.1B tokens
COEM-Large (768d)	25%	\$750K	\$0.12/1M tokens	6.25B tokens
COEM-XL+ (1024d+)	10%	\$300K	\$0.20/1M avg	1.5B tokens
Total Tokens				43.4B/year
Monthly Average				3.6B tokens/month
Managed Infrastructure (30% = \$1.5M/year)				
1-2 vCPU Instances	30%	\$450K	\$30/mo avg	1,250 instance-mo
4-8 vCPU Instances	35%	\$525K	\$300/mo avg	1,750 instance-mo
16+ vCPU Clusters	25%	\$375K	\$3,500/mo avg	107 cluster-mo
Multi-region (Giants)	10%	\$150K	\$50K/mo avg	3 cluster-mo
Avg Monthly Infrastructure				\$125K/month
Professional Services (10% = \$500K/year)				
Enterprise Integrations	50%	\$250K	\$75K avg	3-4 projects
Custom Model Training	30%	\$150K	\$100K avg	1-2 projects
Onboarding & Training	20%	\$100K	\$15K avg	6-7 engagements

Table 9: Year 2 Sales Targets by Revenue Stream

Year 2 Customer Mix (400 customers total):

- **200 Hobbyist/Starter** @ \$42/mo avg = \$100K ARR (2%)
- **120 Small Startups** @ \$556/mo avg = \$800K ARR (16%)
- **60 Medium Companies** @ \$2,083/mo avg = \$1.5M ARR (30%)
- **15 Large Enterprises** @ \$10,000/mo avg = \$1.8M ARR (36%)
- **5 Tech Giants** @ \$13,333/mo avg = \$800K ARR (16%)

13.1.3 Year 3: \$18.7M ARR Target

Revenue Stream	Mix	Target	Unit Price	Units Required
Model Usage Fees (55% = \$10.3M/year)				
COEM-Small (256d)	20%	\$2.06M	\$0.04/1M tokens	51.5B tokens
COEM-Base (512d)	30%	\$3.09M	\$0.08/1M tokens	38.6B tokens
COEM-Large (768d)	30%	\$3.09M	\$0.12/1M tokens	25.75B tokens
COEM-XL+ (1024d+)	20%	\$2.06M	\$0.22/1M avg	9.4B tokens
Total Tokens				125.3B/year
Monthly Average				10.4B tokens/month
Managed Infrastructure (35% = \$6.55M/year)				
1-2 vCPU Instances	15%	\$982.5K	\$30/mo avg	2,729 instance-mo
4-8 vCPU Instances	25%	\$1.64M	\$300/mo avg	5,458 instance-mo
16+ vCPU Clusters	35%	\$2.29M	\$3,500/mo avg	654 cluster-mo
Multi-region (Giants)	25%	\$1.64M	\$75K/mo avg	22 cluster-mo
Avg Monthly Infrastructure				\$546K/month
Professional Services (10% = \$1.87M/year)				
Enterprise Integrations	40%	\$748K	\$100K avg	7-8 projects
Custom Model Training	35%	\$655K	\$150K avg	4-5 projects
Onboarding & Training	25%	\$468K	\$20K avg	23 engagements

Table 10: Year 3 Sales Targets by Revenue Stream

Year 3 Customer Mix (1,500 customers total):

- **800 Hobbyist/Starter** @ \$31/mo avg = \$300K ARR (1.6%)
- **450 Small Startups** @ \$370/mo avg = \$2.0M ARR (10.7%)
- **180 Medium Companies** @ \$2,315/mo avg = \$5.0M ARR (26.7%)
- **60 Large Enterprises** @ \$10,000/mo avg = \$7.2M ARR (38.5%)
- **10 Tech Giants** @ \$35,000/mo avg = \$4.2M ARR (22.5%)

13.2 Embedding Volume Breakdown

Converting token targets to practical embedding operations:

Metric	Year 1	Year 2	Year 3
Total Tokens/Year	3.8B	43.4B	125.3B
Tokens/Month	317M	3.6B	10.4B
Avg tokens/embedding	150	150	150
Embeddings/Year	25.3M	289M	835M
Embeddings/Month	2.1M	24M	70M
Embeddings/Day	69K	792K	2.3M

Table 11: Embedding Volume Requirements (assuming 150 tokens/embedding average)

13.3 Monthly Sales Velocity Requirements

Year 1 (Months 6-12, post-launch):

Segment	New Customers/Mo	MRR/Customer	Monthly MRR Add
Hobbyist/Starter	4-5	\$55	\$220-275
Small Startups	2-3	\$750	\$1,500-2,250
Medium Companies	1	\$1,250	\$1,250
Enterprise Pilots	0.3 (1 per quarter)	\$1,250	\$375
Total Monthly	7-9		\$3,345-4,150

Table 12: Year 1 Monthly Sales Velocity (Design Partner Phase)

Key insight: Year 1 is about **validating product-market fit**, not maximizing revenue. Focus on converting 5-10 design partners to paid customers and establishing repeatable sales motion for Year 2.

13.4 Gross Margin by Revenue Stream

Revenue Stream	Gross Margin	Y3 Revenue	Rationale
COEM Model Usage	80-90%	\$10.3M	GPU inference cost is \$0.01-0.02/1M tokens; we charge \$0.04-0.30
Managed Infrastructure	0%	\$6.55M	At-cost pass-through (competitive moat)
Professional Services	40-50%	\$1.87M	Labor-intensive but high-ACV
Blended Gross Margin	68%	\$18.7M	Model revenue dominates margin

Table 13: Gross Margin by Revenue Stream (Year 3)

Investor insight: Despite zero-margin infrastructure, our blended gross margin is 68% because **high-margin COEM model usage** (80-90% margin) represents 55% of revenue. This is the inverse of competitors who mark up infrastructure but commoditize embeddings.

13.5 Revenue Sensitivity Analysis

If we miss enterprise sales:

- Year 1 is not dependent on enterprise—designed around design partner conversions
- Year 2: -1 enterprise customer = -\$120K ARR (2.4% reduction)
- Mitigation: SMB self-service provides revenue floor

If token pricing pressure increases:

- **-25% token pricing** = requires 33% more volume to maintain revenue
- Year 3: 167B tokens (vs. 125.3B baseline)
- Mitigation: COEM efficiency improvements and higher-dimension upsells

Upside scenarios:

- **Faster SMB adoption** = 30 small startups in Year 1 = +\$90K ARR
- **Early tech giant pilot** in Year 1 = +\$100-200K ARR

14 INTELLECTUAL PROPERTY

14.1 Proprietary Technology

- **CVD Search Algorithm:** Trade secret - deterministic vector traversal method
- **CosmaDistance Metric:** Proprietary distance formulation optimized for cross-lingual search
- **COEM Training Pipeline:** Trade secret - alignment methodology and loss functions
- **Quantization Methods:** Trade secret - vector compression with minimal quality loss

14.2 Licensing Strategy

- Closed-source core technology
- Open-source client SDKs (community building)
- API-first distribution (SaaS default)
- Enterprise source licensing (restricted redistribution)

15 RISK FACTORS & MITIGATION

As with any early-stage technology company, Azetta.ai faces several categories of risk. We have identified and assessed these risks and developed mitigation strategies for each.

15.1 Market & Competitive Risks

15.1.1 Risk: Competitive Replication

Description: Competitors or well-funded entrants could attempt to replicate our technology, potentially commoditizing our offerings.

Mitigation:

- **Trade Secret Strategy:** Our core innovations (CVD, COEM, CSL) are maintained as proprietary trade secrets, never published or patented. This prevents disclosure while main-

taining exclusive use.

- **API-Only Access:** Technology is accessible only through our API—no open-source components, no model downloads, no architectural disclosure
- **Zero Build Time Moat:** CVD's instant document insertion (vs HNSW's hours-long index rebuilds) is an architectural advantage competitors cannot easily replicate without ground-up redesign
- **Operational Moat:** The complexity of our deployment, optimization, and operational expertise creates a significant barrier to replication. Even if competitors understand the concept, execution at scale requires years of specialized knowledge.
- **Implementation Complexity:** System requires deep expertise across vector databases, multilingual NLP, distributed systems, and kernel methods
- **First-Mover Advantage:** Building deep integration partnerships and customer lock-in while others are still developing
- **Network Effects:** Our multi-tenancy architecture creates data advantages that improve with scale
- **Continuous Innovation:** Ongoing R&D keeps us ahead of potential copycats—by the time someone replicates our current system, we're already two generations ahead

15.1.2 Risk: Big Tech Competition

Description: Google, Microsoft, or Amazon could enter the omnilingual semantic search space with significant resources.

Mitigation:

- **Specialized Focus:** Unlike general-purpose AI providers, we are laser-focused on omnilingual search optimization
- **Cost Efficiency:** Our architecture is fundamentally more cost-efficient than transformer-based approaches at scale
- **Partnership Potential:** We complement rather than compete with cloud providers; potential acquisition target
- **Nimble Execution:** As a startup, we can iterate and deploy faster than large organizations
- **Customer Intimacy:** Deep integration with early customers creates switching costs

15.1.3 Risk: Market Timing

Description: The market for omnilingual semantic search may not develop as quickly as projected.

Mitigation:

- **Early Customer Validation:** Focusing on design partnerships and LOIs to validate demand
- **Horizontal Applicability:** Our technology works across industries, reducing dependence on any single vertical
- **Cost Savings Appeal:** Even without growth, our 60% cost reduction drives adoption in existing search workloads
- **Pivot Flexibility:** Core technology applicable to other semantic understanding use cases

15.2 Technical Risks

15.2.1 Risk: Scaling Challenges

Description: Performance or reliability issues could emerge at enterprise scale.

Mitigation:

- **Proven Architecture:** CVD and COEM designs tested on large-scale datasets
- **Gradual Rollout:** Phased deployment starting with smaller customers
- **Multi-Cloud Strategy:** Redundancy across AWS, GCP, and Azure
- **Observability:** Comprehensive monitoring and alerting infrastructure from day one
- **Expert Team:** Deep distributed systems and ML infrastructure expertise on founding team

15.2.2 Risk: Model Performance

Description: Our models may not deliver expected accuracy or language coverage in production.

Mitigation:

- **Benchmark Validation:** Extensive testing on standard datasets before production
- **Continuous Improvement:** Active learning pipelines to improve performance over time
- **Hybrid Approach:** Fallback to established embedding models when needed
- **Transparent Metrics:** Clear communication of performance tradeoffs with customers

15.3 Execution Risks

15.3.1 Risk: Customer Acquisition

Description: Longer-than-expected sales cycles or high CAC could impact runway.

Mitigation:

- **PLG Motion:** Self-serve tier reduces dependency on enterprise sales

- **Design Partners:** Co-development approach with early customers
- **Content Marketing:** Technical thought leadership to drive inbound interest
- **Partner Channels:** Leverage system integrators and cloud marketplaces
- **Capital Efficient:** 18-month runway provides time to iterate go-to-market

15.3.2 Risk: Operational Complexity

Description: Managing multi-cloud infrastructure and customer deployments is operationally intensive.

Mitigation:

- **Automation First:** Heavy investment in IaC and deployment automation
- **Standard Configurations:** Limited SKUs reduce operational burden
- **Managed Services Tier:** SaaS option for customers who want hands-off
- **Observability:** Proactive monitoring reduces reactive firefighting

15.4 Regulatory & Compliance Risks

15.4.1 Risk: Data Privacy Regulations

Description: GDPR, CCPA, and emerging AI regulations could impact operations.

Mitigation:

- **Privacy by Design:** Data minimization and encryption built in from day one
- **Regional Deployment:** Data residency options for EU, US, and APAC
- **Compliance Roadmap:** SOC 2 Type II and ISO 27001 on 12-month roadmap
- **Customer Control:** Customers retain full control over their data

15.4.2 Risk: AI Liability & Ethics

Description: Emerging regulations around AI bias, explainability, and liability.

Mitigation:

- **Transparent Models:** CVD design is more interpretable than blackbox transformers
- **Bias Testing:** Regular fairness audits across languages and demographics
- **Customer Responsibility:** Clear ToS defining customer responsibility for use cases

15.5 Financial Risks

15.5.1 Risk: Longer Runway Needed

Description: Achieving product-market fit may take longer than 18 months, requiring additional capital.

Mitigation:

- **Milestone-Based Approach:** Clear go/no-go decisions at 6, 12, and 18 months
- **Capital Efficiency:** cloud-based development minimize burn
- **Revenue Acceleration:** Self-serve PLG tier can generate early revenue
- **Strong Investor Network:** Founders have access to follow-on capital if needed
- **Bridge Options:** Strategic partnerships could provide non-dilutive funding

15.6 Risk Management Philosophy

Our approach to risk management is proactive and transparent:

1. **Regular Assessment:** Monthly risk reviews with board/advisors
2. **Scenario Planning:** Best case, base case, worst case planning for all key decisions
3. **Transparent Communication:** Honest dialogue with investors about challenges
4. **Rapid Iteration:** Quick pivots when data suggests strategy adjustment needed
5. **Focus on What We Can Control:** Execution excellence and customer success over market speculation

While these risks are real, our team's experience, technical differentiation, and capital-efficient approach position us well to navigate the challenges of building a category-defining company in the semantic search space.

16 THE ASK: \$1.59M PRE-SEED ROUND

16.1 Funding Request

Azetta.ai is raising **\$1.59 million** in pre-seed funding to build and validate Cosma as the leading omnilingual semantic search platform. This capital will provide an **18-month runway** to achieve product-market fit and position the company for a successful Series A.

Why \$1.59M? Deep-tech AI startups face unique risks: GPU pricing volatility, extended enterprise sales cycles, and the need to attract senior ML talent in a competitive market. This raise provides the buffer to execute properly rather than running lean to the point of fragility.

16.2 Use of Funds

The funding allocation balances aggressive product development with realistic operational needs for a deep-tech AI company:

Engineering Equity Note: Below-market cash compensation for engineering hires (\$160K vs. \$200K+ market rate) is offset by significant equity grants (1.0%-2.0% for Senior ML Engineer, 0.5%-1.0% for Backend Engineer). This aligns incentives and extends runway while attracting talent motivated by ownership.

16.3 Key Milestones with This Funding

The \$1.5M will enable Azetta.ai to achieve the following critical milestones over 18 months:

16.3.1 Technical Milestones (Months 1-9)

- **Q1 (Months 1-3):** Production-ready CVD v1.0 with deterministic sub-50ms retrieval across 10M+ vectors
- **Q2 (Months 4-6):** COEM multilingual models trained on 50+ languages with SOTA cross-lingual transfer performance
- **Q3 (Months 7-9):** Cosma Search Layer (CSL) deployed with auto-scaling, multi-tenancy, and enterprise security (SOC 2 Type I)

16.3.2 Commercial Milestones (Months 6-18)

- **Q2-Q3 (Months 6-9):** 5-10 design partnerships across target verticals (e-commerce, SaaS, finance); establish product-market fit signals
- **Q4 (Months 10-12):** Convert design partners to paying customers (5-10 customers); achieve \$300K-350K ARR from design partner conversions
- **Month 12 (End of Year 1):** Reach **\$350K ARR** with 50-80 total customers (predominantly startups and SMBs, with 1-2 enterprise pilots)
- **Months 13-18:** Scale to **\$1.5M-2M ARR** run rate with 200-400 customers; begin enterprise sales motion with 6-9 month pipeline

Why We're Conservative on Year 1: Enterprise sales cycles are 9-15 months. Our Year 1 revenue comes from design partner conversions and SMB self-service—not enterprise deals. Enterprise ARR materializes in Year 2 once the pipeline matures.

These milestones are ambitious but achievable because of how we've architected both our product *and* our business model for capital efficiency.

16.4 Capital Efficiency Strategy

Azetta.ai is optimizing burn rate through:

- **Lean Core Team:** Founders take below-market salaries (\$80K/year vs. \$150K+ market rate) to extend runway
- **Strategic Compute Spending:** Leverage spot instances and government research grants to reduce GPU costs by 40-60%
- **Design Partner Model:** Early customers provide feedback and validation in exchange for discounted pricing, reducing CAC to near-zero for first 10 customers
- **Open-Source Leverage:** Release CVD core as open-source to drive developer adoption and community validation, reducing marketing costs

Beyond these operational efficiencies, our business model itself is designed to create structural

advantages that reduce risk and accelerate growth.

16.5 Business Model Moats: Built-In Defensibility

Beyond product differentiation, Azetta.ai's pricing architecture creates **structural competitive advantages** that de-risk this investment:

16.5.1 1. Self-Funding Growth Engine

Annual infrastructure commitments mean customers prepay 12 months of cloud costs upfront:

- **Year 1:** \$105K in prepaid cash (30% of \$350K ARR from design partner conversions)
- **Year 2:** \$1.5M in prepaid cash (30% of \$5M ARR)
- **Year 3:** \$5.6M in prepaid cash (30% of \$18.7M ARR)

Cash Flow Assumption: Our 18-month runway is calculated assuming **monthly billing** as the default. Annual prepayments provide upside acceleration but are not required for survival. This is critical: we do not depend on 100% prepayment compliance—enterprise customers typically demand Net-30 or Net-60 terms.

Investor impact: By Year 3, we're generating **\$5.6M+ in upfront cash** annually—sufficient to extend runway and reduce Series B dependency.

16.5.2 2. Competitive Immunity from At-Cost Pass-Through

Our **zero-markup infrastructure pricing** creates an impossible-to-beat moat:

- Competitors charge 20-40% markup on infrastructure (\$6K-7K for workloads we charge \$5K for)
- **They cannot undercut us**—matching our pricing destroys their margins
- We profit from high-margin intelligence (COEM models), not commodity infrastructure
- Enterprise buyers value radical transparency—builds trust and long-term partnerships

Defensive positioning: Competitors must either (1) match our transparency and lose margin, or (2) maintain markups and lose customers. We win either way.

16.5.3 3. Expansion Revenue Built Into Architecture

Zero-fee scaling (enabled by decoupled storage) drives organic expansion:

- Customers upgrade **without penalty**—eliminates friction that kills upsells at competitors
- Target: **110-130% Net Revenue Retention (NRR)** by Year 3
- 20-30% of customers pay usage overage monthly (automatic expansion revenue)
- Can downgrade during downturns without penalty—reduces churn risk to <5%

Math: Every \$1 of new ARR becomes \$1.10-\$1.30 the following year through natural expansion. This compounds: \$1M ARR → \$1.1M → \$1.21M → \$1.33M over 3 years *with zero additional sales effort*.

16.5.4 4. Technical Moat = Pricing Moat

Our decoupled architecture isn't just a technical feature—it's a **pricing differentiator competitors can't match**:

- Traditional vector databases couple storage + compute, forcing customers into rigid instance tiers
- We decouple storage, enabling zero-fee scaling up/down

16.6 Post-Series A Vision (18-36 Months)

With Series A capital (\$5-8M), Azetta.ai will:

1. **Scale GTM:** Build dedicated sales team (3-5 AEs) and expand to 100+ enterprise customers (\$5M ARR)
2. **Product Expansion:** Add vertical-specific features (e-commerce, legal tech) and enterprise integrations (Salesforce, ServiceNow)
3. **Global Expansion:** Establish European presence for GDPR-compliant deployments and multilingual support
4. **R&D Investment:** Expand COEM to 100+ languages and build advanced features (RAG, multi-modal search)

This clear progression—\$1.5M pre-seed to prove product-market fit, Series A to scale GTM, Series B to dominate market—gives investors confidence in a predictable path to a category-defining company.

Bottom line: The \$1.5M pre-seed is not about achieving profitability—it's about proving that Cosma solves a painful, expensive problem better than anyone else, with a clear path to \$5M ARR by Year 2, \$18.7M ARR by Year 3, and positioning for a strong Series B. **Contact Information**

Azetta.ai

This document contains confidential and proprietary information. Distribution is restricted to qualified investors under NDA.

A COMPETITOR PRICING DETAILS

A.1 Voyage AI Detailed Pricing

A.1.1 Text Embedding Models

A.1.2 Multimodal Embeddings

- **Model:** voyage-multimodal-3
- **Text Price:** \$0.12 per 1M tokens
- **Image Price:** \$0.60 per 1B pixels
- **Free Tier:** 200M text tokens + 150B pixels
- **Image Cost Constraints:** Min \$0.00003 (small images), Max \$0.0012 (large images)

A.1.3 Rerankers

A.2 Vector Scaling Costs

- **Voyage AI:**
 - *Scaling Model:* No direct storage costs (embedding provider only).
 - *Cost Factor:* Costs scale linearly with embedding generation volume (tokens processed).
 - *Optimization:* Offers Matryoshka (256-2048 dims) and quantization (int8/binary) to reduce downstream storage costs by up to 6x.
- **Weaviate (Serverless):**
 - *Scaling Model:* Charged per 1 million stored dimensions per month.
 - *Standard Tier:* ≈\$0.095 per 1M dimensions/month (starts at \$25/mo).
 - *Business Critical:* ≈\$0.175 per 1M dimensions/month.
 - *Calculation:* Total Cost = (Number of Vectors × Dimensions per Vector) × Rate.
- **Pinecone (Serverless):**
 - *Scaling Model:* Charged per GB of stored data per month.
 - *Storage Rate:* \$0.33 per GB/month.
 - *Includes:* Vector data (dense/sparse), metadata, and IDs.
 - *Note:* Separate costs for Read/Write units (\$8.25 per 1M units).
 - **Milvus / Zilliz Cloud (Serverless):**
 - * *Scaling Model:* Consumption-based (Virtual Compute Units - vCUs).
 - * *Pricing:* Starts at ≈\$4 per 1M vCUs (Serverless).
 - * *Complexity:* Costs are opaque; vCUs are burned for both background storage maintenance and active compute/querying.
 - * *Storage:* Separate active storage fees apply, often requiring manual configuration of hot/cold tiering to optimize.
 - **pgvector (Self-Hosted/AWS RDS):**
 - * *Scaling Model:* Infrastructure-based (Storage + RAM).
 - * *Memory Requirement:* High. HNSW index typically requires 2-3x the raw vector size in RAM for optimal performance.
 - * *Scaling Cost:* Step-function cost increases when upgrading to larger RAM instances (e.g., AWS r6g.xlarge to 2xlarge).

- * *Optimization*: `halfvec` (2-byte floats) and quantization can reduce RAM needs by 50%+.

A.3 Search Cost Scaling (1M vs 100M vs 1B Vectors)

– Voyage AI:

- * *Search Cost*: **\$0** (External). Voyage charges only for embedding the query text.
- * *Scaling*: Constant cost per query, regardless of database size (1M or 1B).

– Weaviate (Serverless):

- * *1M Vectors*: Covered by Standard Tier ($\approx$ \$25/mo).
- * *100M Vectors*: Cost is driven by **storage** (stored dimensions), not query volume. High monthly storage fee ($\approx$ \$1,500+ for 1536d), but unlimited queries included.
- * *1B Vectors*: Requires Enterprise/Business Critical tier. Primary cost driver is the massive storage footprint, not the search operation itself.

– Pinecone (Serverless):

- * *Query Cost*: Measured in Read Units (RUs).
- * *Scaling Behavior*: **Sublinear**. Querying a 4x larger index does *not* cost 4x RUs (e.g., 50k vectors $\approx$ 5 RUs, 200k vectors $\approx$ 8 RUs).
- * *1B Vectors*: While per-query RU efficiency improves, the sheer volume of data requires dedicated read nodes or high shard counts to maintain latency, increasing base infrastructure costs.
- * **Milvus / Zilliz Cloud**:
 - *1M Vectors*: Free tier covers basic dev/test usages (up to 2.5M vCU/mo).
 - *100M Vectors*: **Complex Tiering**. Requires manual management of "Hot" (RAM/SSD) vs "Cold" (S3) collections to control costs. Active high-performance indexes burn vCUs continuously even when idle.
 - *1B Vectors*: **High OpEx Volatility**. While cheaper than pure RAM (HNSW), the "Serverless" vCU burn creates unpredictable month-to-month bills compared to fixed provisioned capacity.
- * **pgvector (Self-Hosted)**:
 - *1M Vectors*: Low cost. Fits in standard RAM (e.g., 4GB instance).
 - *100M Vectors*: **The RAM Cliff**. HNSW index for 1536d vectors requires $\approx$ 600GB RAM. Requires expensive high-memory instances (e.g., AWS r6g.24xlarge).
 - *1B Vectors*: **Prohibitive w/o Quantization**. Uncompressed HNSW requires $\approx$ 6TB RAM. Practical implementation mandates binary quantization or disk-based indexing to avoid \$20k+/month infrastructure bills.

Category	Item	Amount
Personnel (55% of budget)		
Core Team Salaries	CEO (Taha Bouhsine) — 18 months	\$150,000
	CTO (Douglas Seo) — 18 months	\$150,000
	COO (Poppy Yuan) — 18 months	\$150,000
Engineering Hires	Senior ML Engineer (salary & benefits) — 18 months	\$240,000
	Backend Engineer (salary & benefits) — 15 months	\$180,000
Subtotal		\$870,000
Research & Compute (23% of budget)		
Model Training	Large-scale multilingual model training (cloud GPU clusters)	\$250,000
Production Infrastructure	API hosting, inference serving, and monitoring	\$50,000
Data Acquisition	Multilingual training datasets and pre-processing	\$15,000
Data Pipelines	Vector generation and search optimization	\$30,000
Subtotal		\$345,000
General & Administrative (18% of budget)		
Office & Operations	Co-working space (\$3K/month × 18 months)	\$54,000
Professional Services	Legal (incorporation, IP), accounting, compliance	\$50,000
Software & Tooling	GitHub, Slack, monitoring, analytics, productivity tools	\$21,000
Marketing & DevRel	Developer relations, content, conferences, community	\$80,000
Contingency	Reserve for GPU volatility, legal issues, delays	\$120,000
Subtotal		\$325,000
Buffer	Pre-Series A bridge / unexpected opportunities	\$50,000
Total Funding Request		\$1,590,000

Model	Price (per 1M tokens)	Free Tier
voyage-3-large	\$0.18	200M tokens
voyage-context-3	\$0.18	200M tokens
voyage-3.5	\$0.06	200M tokens
voyage-3.5-lite	\$0.02	200M tokens
voyage-code-3	\$0.18	200M tokens
voyage-finance-2	\$0.12	50M tokens
voyage-law-2	\$0.12	50M tokens
voyage-code-2	\$0.12	50M tokens

Table 14: Voyage AI Text Embedding Pricing

Model	Price (per 1M tokens)	Free Tier
rerank-2.5	\$0.05	200M tokens
rerank-2.5-lite	\$0.02	200M tokens

Table 15: Voyage AI Reranker Pricing