

Cosma

Omnilingual Semantic Search Platform

Technical Whitepaper - Investor Edition

Azetta.ai

Version 1.0 - January 2, 2026

Confidential & Proprietary

Do not distribute without permission.

Executive Summary

Azetta.ai delivers **Cosma**, an enterprise-grade omnilingual semantic search platform built on proprietary vector database technology and unified multilingual embeddings. Our platform enables organizations to search across 100+ languages with a single query, deployed as a fully managed SaaS solution.

The Problem: AI infrastructure is approaching an economic cliff. Enterprises face two fatal options:

- **The Filter Trap:** Partition data by language to control costs, sacrificing global intelligence
- **The RAM Cliff:** Attempt pure search requiring 1.5-3TB RAM for 1B vectors (\$15k+/month)

The Solution: Cosma eliminates this trade-off through three proprietary innovations:

1. **CVD:** A “Structure-Free” vector database that runs on disk with O(1) memory overhead
2. **COEM:** Omnilingual embeddings mapping 100+ languages into a unified vector space
3. **CSL:** Production-ready search layer with sub-50ms latency and 100% recall

Why Now: The **Omnilingual Age** is inevitable. Language barriers in information retrieval are collapsing. We are not optimizing existing infrastructure—we are building the foundation that makes global semantic search profitable.

The Ask: \$1.59M pre-seed for 18-month runway to achieve \$350K ARR (Year 1) and position for Series A at \$5M+ ARR (Year 2).

Contents

1 MARKET THESIS: THE OMNILINGUAL AGE

We stand at the precipice of a new era in human knowledge: the **Omnilingual Age**. The artificial barriers that have divided information for millennia—language, script, and geography—are collapsing. In this new world, a query in English must instantly unlock insights in Mandarin, Arabic, or Hindi.

This is not a feature request; it is an inevitability. But the infrastructure to support it does not exist.

1.1 The Infrastructure Crisis

The AI infrastructure market is approaching a critical **economic cliff**. As enterprises transition from building prototypes to deploying global-scale applications, they are discovering that the current generation of vector databases is mathematically incapable of scaling without destroying unit economics.

The Core Dilemma: Scale vs. Scope

Today's architects are forced to choose between two fatal options:

1. **Limit the Scope (The Filter Trap)**: Restrict search to specific languages or metadata tags to keep costs low, but sacrifice global intelligence.
2. **Pay the Scale Tax (The RAM Cliff)**: Attempt "pure" search over the entire dataset, requiring massive RAM usage that makes unit economics impossible.

1.1.1 The Filter Trap: Artificial Intelligence in Silos

To survive existing infrastructure costs, 95% of enterprise deployments rely on **language-specific filters**. They artificially partition data—isolating "English Customer Support" from "Spanish Technical Docs"—to reduce the search space.

This is a functional regression. It re-imposes the very silos that vector search was promised to destroy. A query in Hindi cannot leverage knowledge from an English manual because the system is explicitly told not to look there. The cost is lowered, but the **intelligence is capped**.

1.1.2 The RAM Cliff: The Economics of "Pure" Search

"Pure" vector search—scanning the high-dimensional manifold to find the absolute best match regardless of partition—is the holy grail of retrieval. However, it requires an infrastructure that does not exist today.

Existing solutions (HNSW, IVFPQ) suffer from a crippling mathematical flaw: **Linear Cost Scaling** ($O(N)$).

- **The RAM Cliff**: To deliver sub-second latency, indices must be held in RAM. As data grows to billions of vectors, memory costs scale non-linearly, forcing companies to buy expensive, high-memory GPU instances just to store the index.
- **The Accuracy Tax**: To avoid this cost, systems use aggressive compression (quantization), which degrades result quality. You pay less, but your AI gets dumber.

1.2 The Jevons Paradox: Unlocking Dormant Demand

The **Jevons Paradox** states that increasing the efficiency of a resource leads to an explosion in its consumption. Today, semantic search is the world's most valuable resource, yet it remains artificially scarce. It is capped by a broken legacy stack that imposes a "Multilingual Tax" on every query and a "Vector Tax" on every byte of data.

Cosma is the catalyst designed to shatter these caps. We are not merely building a faster database; we are engineering the efficiency breakthrough that will trigger the Jevons Paradox for global intelligence.

1.3 Stifled Innovation: The Unbuilt Applications

This infrastructure gap is not just an inconvenience; it effectively **bans** the creation of high-value applications. A class of transformative products remains theoretically possible but economically unbuildable:

- **Global Information Engines:** Systems that instantly synthesize insights from a corporation's *total* knowledge base, across 100+ languages, without manual routing.
- **Planetary-Scale Biometrics:** Face recognition authentication that functions seamlessly across billions of identities, vital for global security and payments.
- **Universal Genomic Matching:** Searching billions of genetic sequences to identify rare disease markers or patient matches without pre-filtering by population or study type.
- **Real-Time Content Integrity:** Instantly checking every frame of user-uploaded video against a global media registry for copyright or safety violations.
- **Unbounded Pattern Discovery:** Fraud detection systems that find subtle correlations across petabytes of unstructured logs without needing pre-defined "suspect" filters.

1.4 The Cosma Monopoly

Cosma is the world's first vertically integrated **Omnilingual Semantic Search Platform**. We have rebuilt the stack from first principles to create a defensible monopoly on the future of search:

1. **Unified Intelligence (COEM):** One vector space for all languages. No translation, no silos.
2. **Infinite Scale (CVD):** A proprietary disk-native engine that turns storage into search, eliminating the RAM bottleneck.
3. **Frictionless Access (CSL):** A production-ready layer that democratizes this power for every developer.

We are not competing for a slice of the existing search market. We are unlocking a new, exponentially larger market—one where the cost of intelligence falls to zero, and the volume of consumption scales to infinity.

The Opportunity

The market does not need another incremental optimization. It needs a **first-principles infrastructure overhaul**. We provide the foundation that eliminates this “Filter vs. Scale” trade-off, enabling pure, global vector search to become the profitable standard for the industry.

1.5 Competitive Positioning

1.5.1 Key Differentiators

1. **True Omnilingual Search:** Single unified index, not language-specific routing
2. **Deterministic Performance:** No ANN approximations = predictable latency
3. **Enterprise Security:** SOC 2 Type II compliance and dedicated isolated environments
4. **Cost Efficiency:** CPU-optimized, 60% lower infrastructure costs vs. GPU-dependent solutions

1.5.2 Target Markets

- **Global Enterprises:** Companies with multilingual content (legal, customer support, documentation)
- **E-commerce:** Cross-border product search and recommendations
- **Healthcare:** Medical literature search across languages
- **Government:** Intelligence, diplomatic communications, public records
- **Legal:** Contract search, due diligence, compliance

Margin Note: Our at-cost infrastructure pass-through creates competitive immunity on price. However, our real margin comes from **COEM model usage fees (80-90% gross margin)**, not infrastructure. This is the inverse of competitors who mark up infrastructure but commoditize embeddings.

2 TECHNICAL CORE: THE “STRUCTURE-FREE” STACK

Cosma is built on a vertically integrated technology stack designed from first principles to eliminate the fundamental trade-offs plaguing existing vector search infrastructure. This section details our three proprietary components—CVD, COEM, and CSL—and the intellectual property that protects them.

2.1 System Architecture

Figure ?? illustrates Cosma’s four-layer architecture designed for enterprise deployment.




Figure 1: Cosma System Architecture - Four-Layer Design

2.2 Cosma Vector Database (CVD)

CVD is a proprietary storage engine that eliminates graph-maintenance overhead, enabling **disk-native** vector search with deterministic performance.

2.2.1 The Architecture of Efficiency

Traditional vector databases rely on graph-based indexing (HNSW) or cluster-based approximations (IVF) to overcome the scaling limit. These legacy architectures introduce unavoidable hardware and operational taxes.

What We Eliminated:

- **No Graph Construction Overhead:** Competitors burn 30-50% of CPU cycles “wiring” indices. CVD eliminates this step, enabling instant availability.
- **No “RAM Cliff”:** Graph indices must reside in memory to be performant, forcing customers to rent expensive high-RAM instances (e.g., AWS r6g.24xlarge). CVD is architected for **Disk-Native Efficiency**, running on standard commodity storage without latency degradation.
- **No Re-Indexing Lag:** CVD removes the need for background “build” or “warm-up” phases. Data is searchable the millisecond it is ingested.

2.2.2 Proprietary Search Execution: Direct-Storage Traversal

CVD employs a trade-secret **Direct-Storage Traversal** algorithm. By decoupling semantic precision from RAM-heavy graph structures, we achieve:

- **Predictable Latency:** Removing the probabilistic variance of graph “hops.”
- **Linear Cost Scaling:** Costs grow with storage (cheap), not memory (expensive).

“Unlike graph-based approximations where query pathing changes based on index fragmentation (leading to jitter), Cosma’s search path is architecturally fixed. This delivers consistent latency profiles regardless of data scale.”

2.2.3 Trade Secret: O(1) Memory Overhead

Our proprietary indexing method is **stateless regarding neighbor relationships** at the storage layer. This means:

- **O(1) Memory Overhead:** We do not store edge pointers.
- **Instant Mutability:** Updates and deletes are standard database operations, not graph repairs.

This architectural decision is protected as a trade secret and is never disclosed through patents or publications.

2.3 Cosma Omnilingual Embedding Models (COEM)

COEM models solve the fundamental challenge of cross-lingual semantic representation by mapping 100+ languages into a **unified vector space** where semantically similar content clusters regardless of language.

2.3.1 Model Architecture




Figure 2: COEM Model Architecture - Unified Cross-Lingual Latent Space

2.3.2 Training Methodology

COEM models achieve cross-lingual alignment through:

1. **Parallel Corpus Training:** Millions of aligned text pairs across 100+ languages
2. **Contrastive Learning:** Pushes semantically similar content together, dissimilar apart
3. **Manifold Regularization:** Ensures smooth geometric structure in latent space
4. **Domain-Specific Augmentation:** Fine-tuning for search-optimized representations

2.3.3 Trade Secret: Novel Mercer Kernel

The core innovation enabling COEM's "Structure-Free" property is a **novel Mercer kernel** that eliminates the need for activation functions and layer normalizations in neural networks.

Key Properties:

- **Primal-Form Kernels:** First method to use Mercer kernels in their primal form for neural networks, bypassing traditional dual-form limitations.
- **Anchor-Free Contrastive Learning:** Removes the need for negative mining, simplifying training dynamics.
- **Disk-Compatible Embeddings:** Vectors produced by COEM are mathematically compatible with disk-based indexing—a property that standard embeddings (OpenAI, Cohere, Voyage) do not possess.

IP Protection: The kernel formulation, training pipeline, and loss functions are maintained as trade secrets and are never published or patented.

2.4 Cosma Search Layer (CSL)

CSL provides the production-ready search API that coordinates embedding generation, vector search, and result ranking.

2.4.1 Search Features

- **Automatic Embedding:** Text automatically converted to vectors
- **Metadata Filtering:** Structured filters (date, category, author)
- **Semantic Re-ranking:** Distance + coherence + language priors
- **Result Normalization:** Consistent scoring across languages

2.5 Intellectual Property Strategy

Our core technology is protected through a deliberate **trade secret strategy**:

- **CVD Search Algorithm:** Trade secret - deterministic vector traversal method
- **CosmaDistance Metric:** Proprietary distance formulation optimized for cross-lingual search
- **COEM Training Pipeline:** Trade secret - alignment methodology and loss functions
- **Quantization Methods:** Trade secret - vector compression with minimal quality loss

Licensing Strategy:

- Closed-source core technology
- Open-source client SDKs (community building)
- API-first distribution (SaaS default)
- Enterprise source licensing (restricted redistribution)

API-Only Access: Technology is accessible only through our API—no open-source components, no model downloads, no architectural disclosure. This prevents disclosure while maintaining exclusive use.

3 BUSINESS MODEL

3.1 Overview

Azetta.ai delivers **Cosma**, an enterprise-grade omnilingual semantic search platform operating on a tiered subscription model with usage-based pricing. We target five distinct customer segments with differentiated value propositions and cost structures, generating revenue through three primary streams.

3.2 Revenue Streams

1. **Model Usage Fees:** Consumption-based pricing for COEM embeddings
 - Dimension-based tiered pricing (see Section ??)
 - Bundled in subscription tiers or pay-as-you-go overage
2. **Managed Infrastructure:** At-cost pass-through pricing
 - *Cloud Cost:* Customer pays raw cloud provider rate (e.g., AWS/GCP instance cost) with no markup.
 - *Billing:* Annual commitment preferred.
 - *Flexibility:* Zero-fee instance scaling (upgrade/downgrade) enabled by decoupled storage architecture.
3. **Professional Services:** Integration, training, custom models
 - Typical engagements: \$50K-\$250K
 - Expected revenue mix: 10-15% (Years 1-2), declining to 5% (Year 5)
 - Services include: Custom model training, integration support, enterprise onboarding

3.3 Strategic Business Model Advantages

Azetta.ai's pricing model is architected to create **sustainable competitive moats** while maximizing customer lifetime value:

3.3.1 1. At-Cost Infrastructure Pass-Through

What it means: Customers pay the raw cloud provider cost (AWS/GCP) for compute and storage with **zero markup**.

Why this wins:

- **Radical Transparency:** No hidden margins or surprise infrastructure bills—customers see exactly what they're paying for
- **Competitive Immunity:** Competitors can't undercut us on infrastructure since we charge zero margin
- **Trust Builder:** Enterprise buyers love knowing we're not profiting from their infrastructure spend—builds long-term partnerships
- **Revenue Focus:** We make money on value-add (COEM models, search intelligence), not commodity infrastructure

Customer impact: A company spending \$5,000/month on infrastructure pays exactly \$5,000. With traditional vendors (20-40% markup), the same workload costs \$6,000-\$7,000/month.

3.3.2 2. Annual Commitments = Predictable Revenue

What it means: Infrastructure costs are billed **annually in advance**, not month-to-month.

Why this wins:

- **12-Month Cash Flow:** Every customer provides upfront capital, reducing working capital needs
- **Lower Churn Risk:** Annual contracts create switching friction—customers don't casually cancel mid-year
- **Financial Predictability:** ARR becomes highly predictable, making revenue forecasting accurate for investors
- **Customer Savings:** We pass along AWS annual commitment discounts (10-15% savings vs. on-demand)

Investor impact: By Year 3, 70%+ of revenue is under annual contracts, providing **\$19.7M+ in prepaid cash** to fund operations without relying on external capital.

3.3.3 3. Zero-Fee Scaling Flexibility

What it means: Customers can scale infrastructure up or down **without change fees or penalties**, enabled by our decoupled storage architecture.

Why this wins:

- **Eliminates Upgrade Friction:** Customers don't fear scaling up (no "gotcha" fees), accelerating expansion revenue
- **Technical Moat:** Decoupled architecture is hard to replicate—competitors use fixed instances
- **Expansion Revenue Driver:** Easy scaling drives 20-30% annual upsells through natural usage growth
- **Risk Mitigation:** Customers can downgrade during slowdowns without penalty, reducing churn during economic downturns

Economic impact: Net Revenue Retention (NRR) of 110-130% driven by frictionless expansion—every \$1 of new ARR becomes \$1.10-\$1.30 the following year.

3.3.4 4. Hybrid Pay-As-You-Go + Subscription Model

What it means: Customers buy bundled tiers (predictable base) + pay overage fees for usage spikes (consumption alignment).

Why this wins:

- **Land-and-Expand:** Start with small subscription, grow into higher tiers as usage increases organically
- **Revenue Upside:** Overage fees capture unexpected growth (20-30% of customers exceed base tier monthly)

- **Budget Certainty:** Base subscription is fixed; overage only applies when customers scale, aligning costs with their revenue growth
- **Self-Service Friendly:** Usage-based pricing removes sales friction—customers self-select appropriate tiers

Revenue mix: 60% subscription (predictable) + 30% infrastructure pass-through (annual) + 10% overage/services (expansion) = highly diversified revenue base.

3.3.5 5. Proprietary Ecosystem Lock-in

What it means: The Cosma Vector Database (CVD) is tightly coupled with our model architecture.

Why this wins:

- **Technical Dependency:** CVD's "zero-build" efficiency relies on specific vector properties produced only by Cosma models or our adaptation API.
- **No "Bring Your Own Vector" Leakage:** Customers cannot simply generate vectors with OpenAI and query CVD directly. They *must* use our adaptation layer, ensuring we capture revenue even if they use third-party models.
- **Sticky Ecosystem:** The database and model reinforce each other—using one makes the other more valuable.

3.3.6 6. Volume Expansion (The Jevons Paradox)

What it means: By reducing the unit cost of vector search by 10x, we unlock massive-scale use cases that were previously cost-prohibitive.

Why this wins:

- **New Workloads:** Instead of just indexing "hot" data (recent 10%), customers can now afford to index *everything* (archival data, logs, full history).
- **Competitive Defensibility:** Pass-through pricing + decoupled architecture = nearly impossible for competitors to match without rebuilding their entire platform

3.3.7 7. The "Structure-Free" & "Bridge" Model Moat

What it means: Current vector databases require massive RAM because they must store complex graphs (HNSW) to understand similarity. Azetta has developed the world's first "**Structure-Free**" **Embedding Model**. Our proprietary model embeds data directly into a format that standard storage engines can sort natively.

Why this wins:

- **5% Cost Basis:** This eliminates the need for the graph entirely, allowing us to run at 5% of the cost of Pinecone with mathematical guarantees on retrieval speed.
- **The "Bridge" Model Monopoly:** You cannot use standard Postgres for this because standard embeddings (from OpenAI, Cohere, or Voyage) are mathematically incompatible with disk-based indexing. If you try to sort OpenAI vectors on a disk, you get garbage results.
- **Zero-RAM Search:** Azetta is not just a database wrapper. We are the only provider of the necessary "bridge" model—the COEM—that translates human language into this high-

efficiency format. Customers pay us because we own the only model in the world that makes "zero-RAM" search possible.

3.4 COEM Model Pricing Tiers

COEM models are available in multiple dimensions, with pricing reflecting computational cost. Subscription tiers bundle prepaid COEM tokens with managed infrastructure and unlimited search queries.

Model	Dims	\$/1M Tokens	Rel. Cost	Use Case
COEM-Tiny	128	\$0.02	1x	High-volume, basic similarity
COEM-Small	256	\$0.04	2x	Standard applications
COEM-Base	512	\$0.08	4x	Production workloads
COEM-Large	768	\$0.12	6x	High-accuracy search
COEM-XL	1024	\$0.18	9x	Premium semantic understanding
COEM-XXL	1536	\$0.30	15x	Maximum precision

Table 1: COEM Pricing by Dimension (100+ Languages, Unlimited Searches)

Pricing factors: Higher dimensions = more GPU memory, larger index storage, lower throughput per GPU, but better semantic precision.

Example: The "Professional" tier (\$999/month) includes 5M embeddings (40B tokens using COEM-Base at $\$0.08/1M = \3.20 in compute) plus hosted infrastructure and unlimited searches.

3.5 Customer Segmentation

Segment	Typical Use Case	Infrastructure Tier	Monthly Embeddings	Price Range
Individuals/ Hobbyists	Personal projects, learning, prototyping	Shared CPU	10K-100K	Free/\$9-29
Small Startups	MVP development, early apps with <10K users	1-2 vCPU instance	100K-1M	\$99-299
Medium Companies	Production apps, 10K-100K users	4-8 vCPU instance	1M-10M	\$999-2,999
Large Enterprises	Multiple products, 100K-1M users	Dedicated cluster (16-32 vCPU)	10M-100M	\$9,999-29,999
Tech Giants	Platform-scale, millions of users	Multi-region clusters	100M+	Custom (6 figures)

Table 2: Customer Segment Characteristics (Unlimited Searches on Hosted Infrastructure)

3.6 Unit Economics by Segment

3.6.1 Individuals & Hobbyists

Revenue Model:

- Free Tier: Shared CPU + 10,000 embeddings/month (unlimited searches, \$0 revenue)
- Starter: \$9/month - Shared CPU + 50,000 embeddings (unlimited searches)
- Pro: \$29/month - 1 vCPU + 200,000 embeddings (unlimited searches)
- **Included Model:** COEM-Small (256 dims, \$0.04/1M tokens)
- Overage: \$0.10 per 10K embeddings beyond quota (COEM-Small)

Cost Analysis (per paying user):

- Infrastructure (shared compute): \$2-4/month
- Model inference (GPU, batched): \$1-3/month
- Storage & bandwidth: \$0.50-2/month
- **Total COGS:** \$3.50-9/month
- **Gross margin:** 0-70% (Free tier subsidized by paid users)

Conversion assumptions: 5-10% of free users convert to paid tiers.

3.6.2 Small Startups

Revenue Model:

- Growth: \$99/month - 1 vCPU instance + 500K embeddings (unlimited searches)
- Business: \$299/month - 2 vCPU instance + 2M embeddings (unlimited searches)
- **Included Model:** COEM-Small (256 dims, \$0.04/1M tokens)
- Overage: \$0.10 per 10K embeddings beyond quota (COEM-Small)

Cost Analysis (per customer):

- Infrastructure (1-2 vCPU): \$20-40/month
- Model inference (GPU): \$10-25/month
- Storage & bandwidth: \$5-10/month
- Support overhead: \$5-10/month (email support)
- **Total COGS:** \$40-85/month
- **Gross margin:** 60-75%

Revenue projections:

- Average deal size: \$150-250/month (including overage)
- CAC: \$500-1,000 (self-service, marketing-driven)
- Payback period: 4-6 months
- Annual LTV: \$1,800-3,000 (assuming 12-18 month retention)

3.6.3 Medium Companies

Revenue Model:

- Professional: \$999/month - 4 vCPU instance + 5M embeddings (unlimited searches)
- Premium: \$2,999/month - 8 vCPU instance + 20M embeddings (unlimited searches)
- **Included Model:** COEM-Base (512 dims, \$0.08/1M tokens)
- Overage: \$0.12 per 10K embeddings (COEM-Base)
- Includes: 99.9% SLA, dedicated support, advanced features

Cost Analysis (per customer):

- Infrastructure (4-8 vCPU): \$200-400/month
- Model inference (GPU): \$100-250/month
- Storage & bandwidth: \$50-100/month
- Support & account management: \$100-150/month
- **Total COGS:** \$450-900/month
- **Gross margin:** 65-75%

Revenue projections:

- Average deal size: \$1,500-3,500/month
- CAC: \$5,000-10,000 (direct sales)
- Payback period: 4-6 months
- Annual LTV: \$30,000-50,000

3.6.4 Large Enterprises

Revenue Model:

- Enterprise: \$9,999/month - 16 vCPU cluster + 50M embeddings (unlimited searches)
- Enterprise Plus: \$29,999/mo - 32 vCPU + 200M embeddings (unlimited searches)
- **Included Model:** COEM-Large (768 dims, \$0.12/1M tokens)
- Overage: \$0.15 per 10K embeddings (COEM-Large)
- Includes: Single-tenant, VPC peering, premium support, compliance

Cost Analysis (per customer):

- Dedicated infrastructure (16-32 vCPU): \$2,000-5,000/month
- GPU compute cluster: \$1,500-3,500/month
- Storage & bandwidth: \$500-1,000/month
- Premium support & CSM: \$1,000-2,000/month
- **Total COGS:** \$5,000-11,500/month
- **Gross margin:** 50-70%

Revenue projections:

- Average deal size: \$15,000-35,000/month
- CAC: \$30,000-60,000 (enterprise sales)
- Payback period: 6-12 months
- Annual LTV: \$200,000-450,000
- Contract length: 1-3 years

3.6.5 Technology Giants

Revenue Model:

- Custom pricing: \$100,000-500,000+/month
- **Included Model:** COEM-XL or COEM-XXL (1024-1536 dims, negotiated)
- Overage: Custom volume-based pricing
- Volume discounts, custom SLAs, white-glove service
- Platform integration, co-marketing opportunities

Cost Analysis (per customer):

- Massive dedicated infrastructure: \$30,000-100,000/month
- GPU fleets: \$20,000-80,000/month
- Storage & global CDN: \$5,000-20,000/month
- Dedicated engineering & support: \$10,000-25,000/month
- **Total COGS:** \$65,000-225,000/month
- **Gross margin:** 35-55%

Revenue projections:

- Average deal size: \$150,000-400,000/month
- CAC: \$100,000-250,000 (strategic sales)
- Payback period: 12-18 months
- Annual LTV: \$2M-5M+
- Contract length: 2-5 years

3.7 Revenue Stream Mix

Our revenue model evolves as the business scales, with infrastructure and services revenue growing alongside core model usage:

Revenue Stream	Years 1-2	Years 3-5
Model Usage Fees	79%	55%
Managed Infrastructure	8%	35%
Professional Services	13%	10%

Table 3: Projected Revenue Mix by Stream

Infrastructure revenue grows as customers scale to larger deployment tiers, while professional services remain stable at 10% as improved documentation and self-service tooling reduce dependency.

3.8 5-Year Revenue Projection

Year 1 Reality Check: Enterprise sales cycles are 9-15 months. Year 1 revenue comes from **design partner conversions** and SMB self-service—not enterprise deals. Enterprise ARR materializes in Year 2 once the sales pipeline matures.

Segment	Year 1	Year 2	Year 3	Year 4	Year 5
Individuals (paid)	\$20K	\$100K	\$300K	\$600K	\$1.0M
Small Startups	\$180K	\$800K	\$2.0M	\$4.0M	\$7.0M
Medium Companies	\$120K	\$1.0M	\$5.0M	\$10M	\$18M
Large Enterprises	\$30K	\$1.3M	\$7.2M	\$18M	\$35M
Tech Giants	\$0	\$600K	\$4.2M	\$12M	\$24M
Total ARR	\$350K	\$3.8M	\$18.7M	\$44.6M	\$85M

Table 4: Projected Annual Recurring Revenue by Segment (Conservative)

3.9 Key Growth Assumptions

1. Customer acquisition ramps progressively: 60 → 400 → 1,500 → 3,500 → 6,000 total customers
2. Year 1 is primarily design partner conversions; enterprise sales motion begins Month 10
3. Average revenue per customer grows from \$486/month (Y1) to \$1,180/month (Y5)

4. Gross margin improves from 55% (Y1) to 70% (Y5) through scale efficiencies
5. **Churn rates:** 5%/mo (hobbyists), 3%/mo (SMB), 1%/mo (enterprise), 0.5%/mo (tech giants)
6. **Net Revenue Retention:** 110-130% driven by usage growth and tier upgrades
7. Infrastructure costs decrease 15-20% annually through optimization and better pricing
8. Enterprise sales cycle: 9-15 months (Year 1 enterprise revenue is pilot conversions only)

3.10 Lean Team Strategy: Capital-Efficient Scaling

Unlike traditional SaaS companies that aggressively hire large sales teams, Azetta.ai is designed for **extreme operational leverage** through automation, product-led growth, and technical excellence.

3.10.1 Headcount Projection

Function	Year 1	Year 2	Year 3	Year 4	Year 5
Engineering (Product & Infra)	3	5	8	12	15
Sales & Marketing	1	2	4	6	8
Customer Success & Support	0.5	2	4	6	8
G&A (Finance, Legal, Ops)	0.5	1	2	3	6
Total Headcount	5	10	18	27	37
Revenue per Employee	\$70K	\$500K	\$1.04M	\$1.65M	\$2.3M

Table 5: Projected Headcount & Revenue per Employee

Year 2 Support Scaling: We add a **Technical Support Engineer** in Year 2 to handle the jump from 60 to 400 customers. This brings customer:support ratio to 1:200 (vs. industry standard 1:30)—aggressive but achievable with strong documentation, self-service tooling, and technical buyers who prefer docs over phone support.

3.10.2 How We Stay Lean

1. **Product-Led Growth (PLG):** Self-service onboarding and freemium tiers eliminate need for large SDR/BDR teams. Target: 70%+ customers acquired through product, not outbound sales
2. **Automated Customer Success:** Built-in usage analytics, proactive alerts, and comprehensive documentation reduce support tickets by 60%. One CSM can manage 100+ accounts vs. industry standard of 30-40
3. **Engineering Leverage:** Investing in infrastructure automation (auto-scaling, monitoring, deployment) means 15 engineers in Year 5 can support \$122M ARR workloads that competitors need 50+ engineers to manage
4. **Founder-Led Sales (Early Stage):** Founders handle enterprise deals through Year 2, deferring expensive sales hires until repeatable playbook is proven
5. **Outsourced Non-Core Functions:** Accounting, legal, HR, and recruiting handled via fractional executives and external vendors rather than full-time hires

3.10.3 Benchmark Comparison

Capital Efficiency Impact: By Year 5, Azetta.ai achieves \$2.3M revenue per employee—**8-15x higher than typical SaaS companies**. This operational leverage positions us to reach profitability by Year 4 with less dilution.

Metric	Azetta.ai	Typical SaaS	Advantage
Revenue per Employee (Year 5)	\$3.5M	\$150K-250K	10-20x
Customer Support Ratio	1:100+	1:30-40	3x leverage
Sales Team % of Headcount	23%	40-50%	50% reduction
Eng % of Headcount	43%	25-30%	Product focus

Table 6: Lean Operations vs. Industry Benchmarks

3.11 Path to Profitability

Metric	Year 1	Year 2	Year 3	Year 4	Year 5
Total Revenue	\$350K	\$3.8M	\$18.7M	\$44.6M	\$85M
COGS (45-30%)	-\$158K	-\$550K	-\$5.99M	-\$13.4M	-\$25.5M
Gross Profit	\$192K	\$3.25M	\$12.71M	\$31.2M	\$59.5M
Gross Margin	55%	85%	68%	70%	70%
Sales & Marketing (as % revenue)	-\$245K (70%)	-\$2.5M (65%)	-\$6.55M (35%)	-\$11.15M (25%)	-\$17M (20%)
R&D (as % revenue)	-\$140K (40%)	-\$1.75M (46%)	-\$5.61M (30%)	-\$11.15M (25%)	-\$17M (20%)
G&A (as % revenue)	-\$53K (15%)	-\$600K (16%)	-\$2.24M (12%)	-\$5.35M (12%)	-\$10.2M (12%)
EBITDA	-\$246K	-\$1.6M	-\$1.69M	+\$3.55M	+\$15.3M
EBITDA Margin	-70%	-42%	-9%	+8%	+18%

Table 7: Projected Income Statement & Path to Profitability

Break-even: EBITDA-positive by Year 4 at \$44.6M ARR. Company prioritizes aggressive market share capture in Years 1-3 while steadily improving operational efficiency. By Year 5, achieves best-in-class SaaS metrics with 70% gross margins and 18% EBITDA margins, demonstrating sustainable unit economics at scale.

Key efficiency drivers:

- **S&M optimization:** Product-led growth (PLG) motion reduces CAC by 60% from Year 1 to Year 5; self-service adoption drives 70%+ of new customers by Year 5
- **R&D leverage:** Core platform matures by Year 3; later-stage R&D focuses on incremental features and vertical expansion rather than foundational rebuilds
- **Infrastructure efficiency:** Economies of scale reduce per-customer compute costs 40% through better GPU utilization and AWS volume discounts

Note: This trajectory balances growth and profitability. More aggressive growth scenario (maintaining 30% S&M spend through Year 5) could achieve \$200M+ ARR but delay profitability to Year 6.

3.12 Consolidated Unit Economics

3.12.1 CAC/LTV Density by Segment

3.12.2 Margin Analysis: The Inverse Model

Our margin structure is fundamentally different from competitors:

How 80-90% COEM margin offsets 0% infrastructure margin:

Segment	CAC	Annual LTV	LTV:CAC	Payback	Gross Margin
Hobbyist/Starter	\$50-100	\$200-350	3-4x	2-4 mo	0-70%
Small Startups	\$500-1K	\$1.8K-3K	3-4x	4-6 mo	60-75%
Medium Companies	\$5K-10K	\$30K-50K	5-6x	4-6 mo	65-75%
Large Enterprises	\$30K-60K	\$200K-450K	6-8x	6-12 mo	50-70%
Tech Giants	\$100K-250K	\$2M-5M+	15-20x	12-18 mo	35-55%

Table 8: Consolidated CAC/LTV Analysis by Customer Segment

Revenue Stream	Gross Margin	Y3 Revenue	Rationale
COEM Model Usage	80-90%	\$10.3M	GPU inference cost is \$0.01-0.02/1M tokens; we charge \$0.04-0.30
Managed Infrastructure	0%	\$6.55M	At-cost pass-through (competitive moat)
Professional Services	40-50%	\$1.87M	Labor-intensive but high-ACV
Blended Gross Margin	68%	\$18.7M	Model revenue dominates margin

Table 9: Margin Analysis by Revenue Stream (Year 3)

Despite zero-margin infrastructure pass-through, our blended gross margin is 68% because high-margin COEM model usage (80-90% margin) represents 55% of revenue by Year 3. This is the **inverse** of competitors who mark up infrastructure but commoditize embeddings.

3.12.3 COEM Inference Cost Basis

- **GPU Cost Basis:** \$0.01-0.02 per 1M tokens (batched inference)
- **Our Pricing:** \$0.04-0.30 per 1M tokens (4-15x markup)
- **Result:** 80-90% gross margin on every embedding operation

The at-cost infrastructure strategy creates **competitive immunity**—competitors cannot undercut our infrastructure pricing without destroying their own margins. Meanwhile, we capture value through proprietary model intelligence that they cannot replicate.

4 OPERATIONAL INFRASTRUCTURE & COMPLIANCE

For enterprise buyers, deployment and security are a single procurement hurdle. This section details Cosma's infrastructure options and compliance posture.

4.1 Deployment Models

Cosma supports three deployment tiers designed for different organizational needs.




Figure 3: Cosma Deployment Tiers - Progressive Service Levels

4.2 Enterprise Tier: Single-Tenant Isolation

The Enterprise tier provides complete isolation for organizations with strict security requirements:

4.2.1 VPC Peering Specifications

- **Private Connectivity:** AWS PrivateLink / GCP Private Service Connect
- **Network Isolation:** Dedicated VPC with customer-controlled CIDR blocks
- **No Internet Exposure:** All traffic remains within cloud provider backbone
- **IP Whitelisting:** Customer-defined ingress/egress rules

4.2.2 Single-Tenant Architecture

- **Dedicated Compute:** Isolated VM instances (not shared with other customers)
- **Dedicated Storage:** Separate PostgreSQL cluster with customer-specific encryption keys
- **Dedicated GPU Pool:** Reserved inference capacity for predictable latency
- **Custom Scaling:** Customer-defined auto-scaling policies

4.3 Security & Data Protection

4.3.1 Encryption Standards

- **In Transit:** TLS 1.3 for all API communications
- **At Rest:** AES-256 encryption with customer-managed keys (BYOK) available
- **Key Management:** AWS KMS / GCP Cloud KMS integration

4.3.2 Data Isolation

- **Namespace Separation:** Logical isolation at application layer (Free/Premium tiers)
- **Physical Isolation:** Dedicated infrastructure (Enterprise tier)
- **Privacy Guarantee:** Customer data is **never** used for model training
- **Audit Logging:** Anonymized access logs with configurable retention

4.4 Compliance Certifications

Cosma meets enterprise compliance requirements:

Certification	Status	Tier Availability
SOC 2 Type II	In Progress (Q2 target)	All tiers
GDPR	Compliant	All tiers
HIPAA	Ready (BAA available)	Enterprise tier
ISO 27001	Data centers certified	All tiers

Table 10: Compliance Certification Status

4.5 Omnilingual Data Residency

The “Omnilingual” nature of Cosma introduces unique data residency considerations. A single unified index may contain text from multiple jurisdictions.

4.5.1 GDPR-Compliant Sharding

- **Regional Deployment:** EU-only CVD instances for GDPR-regulated data
- **Data Locality Tags:** Metadata tagging to enforce regional query routing
- **Cross-Region Search:** Federated queries across regional indexes (Enterprise tier)
- **Right to Deletion:** Efficient per-document deletion without index rebuild

4.5.2 Supported Regions

- **US:** AWS us-east-1, us-west-2; GCP us-central1
- **EU:** AWS eu-west-1, eu-central-1; GCP europe-west1
- **APAC:** AWS ap-southeast-1; GCP asia-east1 (roadmap)

5 EXECUTION ROADMAP & RISK MITIGATION

A roadmap without associated risks is a marketing document. This section links each milestone to its specific risks and mitigations, demonstrating engineering maturity.

5.1 Milestone-Risk Matrix

Milestone	Timeline	Key Risk	Mitigation
CVD v1.0 Production	Q1 (Months 1-3)	Scaling challenges at enterprise load	Phased rollout starting with smaller customers
COEM 50+ Languages	Q2 (Months 4-6)	Model performance across low-resource languages	Hybrid fallback to established models
8-bit Quantization	Q2-Q3	Accuracy degradation	Multi-probe search compensation
CSL Enterprise Security	Q3 (Months 7-9)	SOC 2 audit delays	Pre-audit readiness assessment
Distributed Sharding	12-24 months	Consistency at scale	Eventual consistency model with conflict resolution
100B+ Vector Scale	24+ months	Infrastructure costs	Cold/hot tiering with intelligent prefetch

Table 11: Milestone-Risk Linkage Matrix

5.2 Technology Roadmap

5.2.1 Near-Term (6-12 months)

- Public SDK release (Python, JavaScript, Go)
- Expanded language support (120+ languages)
- Index compression (8-bit quantization)
- Real-time indexing pipeline

5.2.2 Mid-Term (12-24 months)

- Distributed sharding for 100B+ vectors
- Streaming search capabilities
- Multi-modal embeddings (text + images)
- Edge deployment (low-power devices)

5.2.3 Long-Term (24+ months)

- Open adapter ecosystem for domain-specific models
- Plugin architecture for custom distance functions
- Federated search across multiple CVD instances
- AutoML for embedding model customization

5.3 Quantization Strategy

5.3.1 8-bit Quantization for COEM

Moving from FP32 to INT8 provides 4x storage reduction with controlled accuracy trade-offs:

- **Storage Reduction:** 4x compression (768 dims × 4 bytes → 768 dims × 1 byte)
- **Expected mAP Impact:** [TODO: Benchmark data needed]
- **Mitigation:** Multi-probe search with re-ranking compensates for quantization error
- **COEM-Large (768d):** Primary candidate for quantization at enterprise scale

5.4 Risk Categories & Mitigations

5.4.1 Market & Competitive Risks

Risk: Competitive Replication

Competitors or well-funded entrants could attempt to replicate our technology.

Mitigation:

- **Trade Secret Strategy:** Core innovations (CVD, COEM, CSL) are proprietary—never published or patented
- **API-Only Access:** No open-source components, no model downloads, no architectural disclosure
- **Zero Build Time Moat:** CVD's instant insertion (vs HNSW's hours-long rebuilds) is an architectural advantage
- **Continuous Innovation:** Ongoing R&D keeps us ahead—by the time someone replicates our current system, we're two generations ahead

Risk: Big Tech Competition

Google, Microsoft, or Amazon could enter the omnilingual semantic search space.

Mitigation:

- **Specialized Focus:** Unlike general-purpose AI providers, we are laser-focused on omnilingual search
- **Cost Efficiency:** Our architecture is fundamentally more cost-efficient at scale
- **Partnership Potential:** We complement cloud providers; potential acquisition target
- **Customer Intimacy:** Deep integration creates switching costs

5.4.2 Technical Risks

Risk: Scaling Challenges

Performance or reliability issues could emerge at enterprise scale.

Mitigation:

- **Proven Architecture:** CVD and COEM designs tested on large-scale datasets
- **Gradual Rollout:** Phased deployment starting with smaller customers
- **Multi-Cloud Strategy:** Redundancy across AWS, GCP, and Azure
- **Observability:** Comprehensive monitoring from day one

Risk: Model Performance

Our models may not deliver expected accuracy across all languages.

Mitigation:

- **Benchmark Validation:** Extensive testing on standard datasets before production
- **Continuous Improvement:** Active learning pipelines
- **Hybrid Approach:** Fallback to established embedding models when needed

5.4.3 Execution Risks**Risk: Customer Acquisition**

Longer-than-expected sales cycles or high CAC could impact runway.

Mitigation:

- **PLG Motion:** Self-serve tier reduces dependency on enterprise sales
- **Design Partners:** Co-development approach with early customers
- **Content Marketing:** Technical thought leadership drives inbound
- **Capital Efficient:** 18-month runway provides iteration time

5.4.4 Regulatory & Compliance Risks**Risk: Data Privacy Regulations**

GDPR, CCPA, and emerging AI regulations could impact operations.

Mitigation:

- **Privacy by Design:** Data minimization and encryption from day one
- **Regional Deployment:** Data residency options for EU, US, and APAC
- **Compliance Roadmap:** SOC 2 Type II and ISO 27001 on 12-month roadmap

5.4.5 Financial Risks**Risk: Longer Runway Needed**

Achieving PMF may take longer than 18 months.

Mitigation:

- **Milestone-Based Approach:** Clear go/no-go decisions at 6, 12, and 18 months
- **Capital Efficiency:** Lean operations minimize burn
- **Revenue Acceleration:** Self-serve PLG tier generates early revenue
- **Strong Investor Network:** Access to follow-on capital if needed

5.5 Risk Management Philosophy

1. **Regular Assessment:** Monthly risk reviews with board/advisors
2. **Scenario Planning:** Best case, base case, worst case for all key decisions
3. **Transparent Communication:** Honest dialogue with investors about challenges
4. **Rapid Iteration:** Quick pivots when data suggests adjustment needed
5. **Focus on Controllables:** Execution excellence over market speculation

6 THE TEAM

6.1 Founding Team

Azetta.ai is led by a technical founding team with deep expertise in AI/ML, full-stack engineering, and scalable systems. The team combines academic research credentials with proven entrepreneurial execution.

6.1.1 Taha Bouhsine - Co-Founder & CEO

Background:

- **Google Developer Expert (GDE)** in AI/ML
- **Education:** M.S. in Electrical & Computer Engineering (Rowan University, 2024), M.S. in Data Science, B.Sc in Math & CS
- **Research Fellowship:** Awarded research fellowship at Rowan University based on research initiated in Morocco, enabling continued collaboration with the Federal Aviation Administration (FAA)
- **Research Focus:** Deep Learning, Multimodality, Representation Learning, Continual Learning, and Self-Supervised Learning
- **Publications:** Published research in high-impact factor journals including *IEEE Transactions on Communications (8 Impact Factor)*, with applications to aviation safety, civil engineering, security, and remote sensing

Government-Funded Research Experience:

- **U.S. Federal Agencies:** Worked on research projects funded by the Department of Defense (DoD), Department of Energy (DoE), and Department of Transportation (DoT)
- **International Collaboration:** CNRST Morocco-funded satellite imagery research
- **Multimodal Expertise:** Extensive experience across all modalities (text, vision, audio, satellite imagery, sensor data)

Previous Founding Experience:

- Founded **MLNomads**, a collaborative ML research company. Built an open-source replacement for Jupyter notebooks that scaled to **2,000 unique users**, demonstrating ability to build and scale developer tools
- Founded **SafetyFirst.ai**, a non-profit research lab focused on creating safe foundational models for robotics

Community Leadership & Thought Leadership:

- **Invited Speaker:** As a recognized Google Developer Expert (GDE), regularly invited by international communities and Google to present on AI/ML expertise
- Founded **GDSC UIZ, Tensorflow User Group** in Agadir (90+ events, 2,000+ attendees)
- Organizer for GDG Agadir and GDG Philly

Key Technical Innovations:

- **Kernel Methods Expert:** Pioneered a novel Mercer kernel that eliminates the need for activation functions and layer normalizations in neural networks
- **Primal-Form Kernels:** Developed the first method to use Mercer kernels in their primal

form for neural networks, bypassing traditional dual-form limitations

- **Anchor-Free Contrastive Learning:** Invented a contrastive learning algorithm that removes the need for negative mining, significantly simplifying training dynamics
- **First-Principles Design:** Designs algorithms from mathematical first principles rather than empirical heuristics, ensuring theoretical soundness

Expertise: Taha's research innovations form the mathematical foundation of Cosma's technology stack. His expertise in kernel methods, representation learning, and cross-lingual embeddings, combined with his first-principles approach to algorithm design, positions him uniquely to lead Azetta's technical vision. His ability to translate theoretical breakthroughs into production systems bridges the gap between academic research and commercial impact.

6.1.2 Poppy Yuan - Co-Founder & COO

Background:

- **Education:** M.S. in Operations Research & Supply Chain Management, Case Western Reserve University (Weatherhead School of Management, 2019)
- **Previous Cofounder & COO:** Ekowsystems Inc. (Transportation Platform, 2022-2024)
- **Operations Expertise:** Strategic operations executive with proven track record in scaling early-stage platforms and building customer success infrastructure

Previous Founding Experience - Ekowsystems Inc.:

- Led transportation platform development from concept to beta stage, managing relationships across an **11-person distributed team**, C-level executives, shareholders, and technology partners
- Built end-to-end customer engagement program serving **250+ beta users** with **70% participation rate** through strategic communication and value delivery
- Managed development budget **exceeding \$100K**, evaluating and implementing multiple SaaS solutions
- Implemented comprehensive customer support infrastructure (Twilio, Zendesk) achieving **95% user satisfaction**
- Established systematic project management framework, **reducing feedback cycles by 40%** and ensuring alignment between business requirements and technical execution
- Created market intelligence system to analyze competitor data, user behavior patterns, and industry trends

Prior Experience:

- **Consulting Analyst, Medical Mutual of Ohio (2019-2022):**
 - Analyzed complex datasets using SQL and Tableau, generating insights that drove **\$2.2M annual revenue growth**
 - Managed relationships with 100+ clients, maintaining **100% delivery rate**
 - Implemented automated workflow systems reducing manual work by 20 hours/week and improving reports by 40%
- **Operations Analyst, Cleveland Kitchen (2018-2019):**
 - Implemented lean inventory management system, achieving **50% reduction in re-work** and **20% improvement in forecasting accuracy**
 - Developed standardized digital systems decreasing errors to below 2% and increasing efficiency by 40%

Expertise: Poppy brings proven operational execution from founding a transportation platform to beta stage with strong user engagement metrics. Her expertise in customer success infrastructure, stakeholder management, and data-driven decision-making positions her to scale Azetta's operations from early adopters to enterprise customers. Her experience evaluating and implementing SaaS tools provides deep understanding of customer needs in the B2B SaaS market. Her M.S. in Operations Research ensures strategic, analytical approaches to scaling challenges.

6.1.3 Douglas Seo - Co-Founder & CTO

Background:

- **Education:** B.S. in Electrical Engineering & Computer Science, UC Berkeley (2022)
- **CS 188-189 (Artificial Intelligence & Machine Learning):** Top 5% of class with Professor recognition
- **Relevant Coursework:** Data Structures & Algorithms, Machine Learning, Distributed Systems, Web Development, Cybersecurity

Proven Technical Leadership & Execution:

- **Workflow (2025) - Founding Engineer & CTO**
 - Built AI-powered workflow automation platform, achieving **3,000+ monthly visits**
 - Tech stack: PostgreSQL, React, TypeScript, Supabase, Prisma ORM
 - Architected scalable backend infrastructure supporting real-time data synchronization
 - <https://workflow.us>
- **Popper LLC (December 2022 - Present) - CEO & Founder**
 - Scaled to **3,000+ active users** with **\$400 MRR** and sustainable unit economics
 - Built production React Native application handling **20,000+ daily Firestore operations**
 - Designed and deployed scalable NoSQL database architecture with real-time sync and offline-first capabilities
 - Implemented production-grade REST APIs on Google Cloud with OAuth authentication, analytics pipelines, and messaging infrastructure
 - Led agile engineering team of **8 developers** through code reviews, sprint planning, and technical roadmap execution
 - Achieved profitability with **60+ business clients** as first-time founder
 - <https://popper.social>

Technical Expertise Relevant to Vector Database Development:

- **Database Architecture:** Hands-on experience designing and scaling both SQL (PostgreSQL) and NoSQL (Firestore) systems for production workloads
- **Distributed Systems:** Built real-time synchronization infrastructure handling thousands of concurrent operations
- **Performance Optimization:** Optimized database queries and indexing strategies for sub-100ms latency requirements
- **Cloud Infrastructure:** Production deployment experience on Google Cloud and Supabase with monitoring, logging, and auto-scaling
- **AI/ML Integration:** Academic foundation in machine learning algorithms and practical

experience integrating NLP/speech models into production systems

Value to Cosma: Douglas brings the rare combination of ML/AI academic rigor from UC Berkeley and battle-tested experience building profitable, scalable applications from scratch. His expertise in database architecture, distributed systems, and full-stack engineering ensures Cosma's vector database can evolve from a prototype to production-grade infrastructure capable of serving millions of semantic search queries with enterprise-level reliability.

Bottom Line: This isn't a team learning on investors' dime. All three co-founders have already built products with thousands of real users, secured institutional validation, and demonstrated the ability to execute under resource constraints. Taha's research innovations solve the core technical challenge, Douglas's engineering expertise ensures production-grade execution, and Poppy's operational excellence guarantees customer success at scale. The combination of deep technical moat, proven execution, and institutional trust makes this an exceptionally strong founding team for a deep-tech pre-seed venture.

7 THE ASK: \$1.59M PRE-SEED ROUND

7.1 Funding Request

Azetta.ai is raising **\$1.59 million** in pre-seed funding to build and validate Cosma as the leading omnilingual semantic search platform. This capital will provide an **18-month runway** to achieve product-market fit and position the company for a successful Series A.

Why \$1.59M? Deep-tech AI startups face unique risks: GPU pricing volatility, extended enterprise sales cycles, and the need to attract senior ML talent in a competitive market.

7.2 Use of Funds

Category	Item	Amount
Personnel (55% of budget)		
Core Team Salaries	CEO, CTO, COO — 18 months	\$450,000
Engineering Hires	Senior ML Engineer + Backend Engineer	\$420,000
Subtotal		\$870,000
Research & Compute (23% of budget)		
Model Training	Large-scale multilingual training	\$250,000
Infrastructure	API hosting, inference, monitoring	\$95,000
Subtotal		\$345,000
General & Administrative (22% of budget)		
Operations	Office, legal, software, marketing	\$325,000
Contingency	GPU volatility, unexpected opportunities	\$50,000
Subtotal		\$375,000
Total Funding Request		\$1,590,000

7.3 Hiring Sequence & Milestone Ownership

The funding request is structured to align engineering hires with specific technical milestones:

Engineering Equity Note: Below-market cash compensation for engineering hires (\$160K vs. \$200K+ market rate) is offset by significant equity grants (1.0%-2.0% for Senior ML Engineer, 0.5%-1.0% for Backend Engineer).

Milestone	Owner	Timeline	Hire Date
CVD v1.0 Production-Ready	Backend Engineer	Q1 (Months 1-3)	Month 1
COEM 50+ Languages	Senior ML Engineer	Q2 (Months 4-6)	Month 1
Real-time Indexing Pipeline	Backend Engineer	Q2-Q3	Month 4
CSL Enterprise Deployment	CTO (Douglas Seo)	Q3 (Months 7-9)	N/A (Founder)
8-bit Quantization	Senior ML Engineer	Q3	Month 6
SOC 2 Type I Compliance	COO (Poppy Yuan)	Q3	N/A (Founder)

Table 12: Hiring Sequence Aligned with Technical Milestones

7.4 Key Milestones

7.4.1 Technical (Months 1-9)

- **Q1:** Production-ready CVD v1.0 with sub-50ms retrieval across 10M+ vectors
- **Q2:** COEM multilingual models trained on 50+ languages
- **Q3:** CSL deployed with auto-scaling, multi-tenancy, and SOC 2 Type I

7.4.2 Commercial (Months 6-18)

- **Months 6-9:** 5-10 design partnerships
- **Month 12:** \$350K ARR with 50-80 customers
- **Month 18:** \$1.5M-2M ARR run rate, positioned for Series A

7.5 Post-Series A Vision (18-36 Months)

With Series A capital (\$5-8M), Azetta.ai will:

1. Scale GTM with dedicated sales team (3-5 AEs)
2. Add vertical-specific features (e-commerce, legal tech)
3. Establish European presence for GDPR deployments
4. Expand COEM to 100+ languages with multi-modal support

The investment thesis: \$1.59M pre-seed to prove product-market fit → \$5M+ ARR by Year 2 → Series A to dominate omnilingual semantic search.

Contact Information

Azetta.ai

This document contains confidential and proprietary information. Distribution is restricted to qualified investors under NDA.

A COMPETITOR PRICING DETAILS

A.1 Voyage AI Detailed Pricing

A.1.1 Text Embedding Models

Model	Price (per 1M tokens)	Free Tier
voyage-3-large	\$0.18	200M tokens
voyage-context-3	\$0.18	200M tokens
voyage-3.5	\$0.06	200M tokens
voyage-3.5-lite	\$0.02	200M tokens
voyage-code-3	\$0.18	200M tokens
voyage-finance-2	\$0.12	50M tokens
voyage-law-2	\$0.12	50M tokens
voyage-code-2	\$0.12	50M tokens

Table 13: Voyage AI Text Embedding Pricing

A.1.2 Multimodal Embeddings

- **Model:** voyage-multimodal-3
- **Text Price:** \$0.12 per 1M tokens
- **Image Price:** \$0.60 per 1B pixels
- **Free Tier:** 200M text tokens + 150B pixels
- **Image Cost Constraints:** Min \$0.00003 (small images), Max \$0.0012 (large images)

A.1.3 Rerankers

Model	Price (per 1M tokens)	Free Tier
rerank-2.5	\$0.05	200M tokens
rerank-2.5-lite	\$0.02	200M tokens

Table 14: Voyage AI Reranker Pricing

A.2 Vector Scaling Costs

- **Voyage AI:**
 - *Scaling Model:* No direct storage costs (embedding provider only).
 - *Cost Factor:* Costs scale linearly with embedding generation volume (tokens processed).
 - *Optimization:* Offers Matryoshka (256-2048 dims) and quantization (int8/binary) to reduce downstream storage costs by up to 6x.
- **Weaviate (Serverless):**
 - *Scaling Model:* Charged per 1 million stored dimensions per month.
 - *Standard Tier:* ≈\$0.095 per 1M dimensions/month (starts at \$25/mo).
 - *Business Critical:* ≈\$0.175 per 1M dimensions/month.
 - *Calculation:* Total Cost = (Number of Vectors × Dimensions per Vector) × Rate.
- **Pinecone (Serverless):**

- *Scaling Model*: Charged per GB of stored data per month.
- *Storage Rate*: \$0.33 per GB/month.
- *Includes*: Vector data (dense/sparse), metadata, and IDs.
- *Note*: Separate costs for Read/Write units (\$8.25 per 1M units).
- **Milvus / Zilliz Cloud (Serverless)**:
 - * *Scaling Model*: Consumption-based (Virtual Compute Units - vCUs).
 - * *Pricing*: Starts at ≈\$4 per 1M vCUs (Serverless).
 - * *Complexity*: Costs are opaque; vCUs are burned for both background storage maintenance and active compute/querying.
 - * *Storage*: Separate active storage fees apply, often requiring manual configuration of hot/cold tiering to optimize.
- **pgvector (Self-Hosted/AWS RDS)**:
 - * *Scaling Model*: Infrastructure-based (Storage + RAM).
 - * *Memory Requirement*: High. HNSW index typically requires 2-3x the raw vector size in RAM for optimal performance.
 - * *Scaling Cost*: Step-function cost increases when upgrading to larger RAM instances (e.g., AWS r6g.xlarge to 2xlarge).
 - * *Optimization*: `halfvec` (2-byte floats) and quantization can reduce RAM needs by 50%+.

A.3 Search Cost Scaling (1M vs 100M vs 1B Vectors)

- **Voyage AI**:
 - * *Search Cost*: **\$0** (External). Voyage charges only for embedding the query text.
 - * *Scaling*: Constant cost per query, regardless of database size (1M or 1B).
- **Weaviate (Serverless)**:
 - * *1M Vectors*: Covered by Standard Tier (≈\$25/mo).
 - * *100M Vectors*: Cost is driven by **storage** (stored dimensions), not query volume. High monthly storage fee (≈\$1,500+ for 1536d), but unlimited queries included.
 - * *1B Vectors*: Requires Enterprise/Business Critical tier. Primary cost driver is the massive storage footprint, not the search operation itself.
- **Pinecone (Serverless)**:
 - * *Query Cost*: Measured in Read Units (RUs).
 - * *Scaling Behavior*: **Sublinear**. Querying a 4x larger index does *not* cost 4x RUs (e.g., 50k vectors ≈ 5 RUs, 200k vectors ≈ 8 RUs).
 - * *1B Vectors*: While per-query RU efficiency improves, the sheer volume of data requires dedicated read nodes or high shard counts to maintain latency, increasing base infrastructure costs.
 - * **Milvus / Zilliz Cloud**:
 - *1M Vectors*: Free tier covers basic dev/test usages (up to 2.5M vCU/mo).
 - *100M Vectors*: **Complex Tiering**. Requires manual management of "Hot" (RAM/SSD) vs "Cold" (S3) collections to control costs. Active high-performance indexes burn vCUs continuously even when idle.
 - *1B Vectors*: **High OpEx Volatility**. While cheaper than pure RAM (HNSW), the "Serverless" vCU burn creates unpredictable month-to-month bills com-

pared to fixed provisioned capacity.

* **pgvector (Self-Hosted):**

- *1M Vectors*: Low cost. Fits in standard RAM (e.g., 4GB instance).
- *100M Vectors*: **The RAM Cliff**. HNSW index for 1536d vectors requires $\approx 600\text{GB}$ RAM. Requires expensive high-memory instances (e.g., AWS r6g.24xlarge).
- *1B Vectors*: **Prohibitive w/o Quantization**. Uncompressed HNSW requires $\approx 6\text{TB}$ RAM. Practical implementation mandates binary quantization or disk-based indexing to avoid \$20k+ /month infrastructure bills.